

INDEX

S.No	Name of the Experiment	Date	Sign
1.	Build an Artificial Neural Network to implement Regression task using the Back-propagation algorithm and test the same using appropriate data sets	15/12/23	
2.	Build an Artificial Neural Network to implement Binary Classification task using the Back-propagation algorithm and test the same using appropriate data sets.	8/01/24	
3.	Build an Artificial Neural Network to implement Multi-Class Classification task using the Back-propagation algorithm and test the same using appropriate data sets.	15/01/24	
4.	Implement an image classification task using pre-trained models like AlexNet, VGGNet, InceptionNet and ResNet and compare the results.	29/01/24	
5.	Design a CNN architecture to implement the image classification task over an image dataset. Perform the Hyper-parameter tuning and record the results.	5/02/24	
6.	Implement a deep neural network (Ex: UNet / SegNet) for Image Segmentation task.	10/02/24	
7.	Train a Recurrent Neural Network (RNN) on the IMDB large movie review dataset for sentiment analysis.	18/02/24	
8.	Download the tweets data from twitter and run an RNN/LSTM to classify the sentiment as Hate/ Non-hate speech. Compare the results of both the models.	4/03/24	
9.	Implement GAN architecture on MNIST dataset to recognize the handwritten digits.	11/03/24	

Experiment:- 1

Aim:- Build an Artificial Neural Network to implement Regression task using the Back-propagation algorithm and test the same using appropriate data sets.

✓ Importing Necessary Libraries

```
import numpy as np
import pandas as pd
```

✓ Loading the Churn Dataset

This data set contains details of a bank's customers and the target variable is a binary variable reflecting the fact whether the customer left the bank (closed his account) or he continues to be a customer.

Binary flag 1 if the customer closed account with bank and 0 if the customer is retained.

```
churn_data = pd.read_csv('Churn_Modelling.csv', delimiter = ',')
churn_data.head(5)
```

0	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
1	2	15647311	Hill	608	Spain	Female	41	1	83807.8
2	3	15619304	Onio	502	France	Female	42	8	159660.8
3	4	15701354	Boni	699	France	Female	39	1	0.0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.8

✓ Accessing the Column Names in the Dataset

```
churn_data.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
       'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
       'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

✓ Setting Column as a Index

```
churn_data = churn_data.set_index('RowNumber')
churn_data.head()
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
1	15634602	Hargrave	619	France	Female	42	2	0.00
2	15647311	Hill	608	Spain	Female	41	1	83807.86
3	15619304	Onio	502	France	Female	42	8	159660.80
4	15701354	Boni	699	France	Female	39	1	0.00
5	15737888	Mitchell	850	Spain	Female	43	2	125510.82

✓ Finding the Shape of the Dataset

```
churn_data.shape
```

```
(10000, 13)
```

```
churn_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 1 to 10000
Data columns (total 13 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   CustomerId        10000 non-null   int64  
 1   Surname           10000 non-null   object  
 2   CreditScore       10000 non-null   int64  
 3   Geography          10000 non-null   object  
 4   Gender             10000 non-null   object  
 5   Age                10000 non-null   int64  
 6   Tenure             10000 non-null   int64  
 7   Balance            10000 non-null   float64 
 8   NumOfProducts      10000 non-null   int64  
 9   HasCrCard          10000 non-null   int64  
 10  IsActiveMember     10000 non-null   int64  
 11  EstimatedSalary    10000 non-null   float64 
 12  Exited             10000 non-null   int64  
dtypes: float64(2), int64(8), object(3)
memory usage: 1.1+ MB
```

✓ Checking Missing Values

```
churn_data.isna().sum()
```

```
CustomerId      0
Surname        0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

✓ Some Columns are Totally Unproductive so let's Remove them

```
churn_data.nunique()
```

```
CustomerId      10000
Surname        2932
CreditScore     460
Geography       3
Gender          2
Age             70
Tenure          11
Balance         6382
NumOfProducts   4
HasCrCard       2
IsActiveMember  2
EstimatedSalary 9999
Exited          2
dtype: int64
```

```
churn_data.drop(['CustomerId','Surname'],axis=1,inplace=True)
```

```
churn_data.head()
```

RowNumber	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
1	619	France	Female	42	2	0.00	1	1	1	101348.88	1
2	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
3	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
4	699	France	Female	39	1	0.00	2	0	0	93826.63	0
5	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
churn_data.shape
```

```
(10000, 11)
```

➤ Some Visualizations

```
[ ] ↓ 7 cells hidden
```

▼ Label Encoding of Categorical Variables

Label Encoding means converting categorical features into numerical values. So that they can be fitted by machine learning models which only take numerical data.

Example: Suppose we have a column Height in some dataset that has elements as Tall, Medium, and short. To convert this categorical column into a numerical column we will apply label encoding to this column. After applying label encoding, the Height column is converted into a numerical column having elements 0,1, and 2 where 0 is the label for tall, 1 is the label for medium, and 2 is the label for short height.

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
churn_data[['Geography', 'Gender']] = churn_data[['Geography', 'Gender']].apply(le.fit_transform)

churn_data.head()
```

RowNumber	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
1	619	0	0	42	2	0.00	1	1	1	101348.88	1
2	608	2	0	41	1	83807.86	1	0	1	112542.58	0
3	502	0	0	42	8	159660.80	3	1	0	113931.57	1
4	699	0	0	39	1	0.00	2	0	0	93826.63	0
5	850	2	0	43	2	125510.82	1	1	1	79084.10	0

▼ Separating Label from Data

```
y = churn_data.Exited
X = churn_data.drop(['Exited'], axis=1)
```

```
X.columns
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
       'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary'],
      dtype='object')
```

```
y
```

```
RowNumber
1      1
2      0
3      1
```

```

4      0
5      0
..
9996   0
9997   0
9998   1
9999   1
10000  0
Name:Exited, Length: 10000, dtype: int64

```

✓ Splitting the Data into Training and Testing

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 2)

print("Shape of the X_train", X_train.shape)
print("Shape of the X_test", X_test.shape)
print("Shape of the y_train", y_train.shape)
print("Shape of the y_test", y_test.shape)

Shape of the X_train (7000, 10)
Shape of the X_test (3000, 10)
Shape of the y_train (7000,)
Shape of the y_test (3000,)

```

Feature Scaling:

The result of **standardization** (or **Z-Score normalization**) is that the features will be re scaled so that they'll have the properties of a standard normal distribution with:

$$\mu = 0$$

And

$$\sigma = 1$$

Where μ is the mean(average) and σ is the standard deviation from the mean; standard scores (also called **Z scores**) of the samples are calculated as follows:

$$z = \frac{x - \mu}{\sigma}$$

About Min-Max Scaling

An alternative approach to **Z-Score** normalization (or called standardization) is the so-called **Min-Max Scaling** (often also simply called **Normalization** - a common cause for ambiguities)

In this approach, the data is scaled to a fixed range - usually $[0, 1]$. The cost of having this bounded range - in contrast to standardization - is that we will end up with smaller standard deviations, which can suppress the effect of outliers.

Note:

If the dataset have lot's of outliers, and the algorithms are sensitive to outliers, please use **Min-Max Scaler**

A Min-Max Scaling is typically done via the following equation:

$$X_{norm} = \frac{X_i - X_{min}}{X_{max} - X_{min}}$$

X_i is the i^{th} sample of dataset.

Z-Score Standardization or Min-Max Scaling

"Standardization or Min-Max scaling"? - There is no obvious answer to this question: it really depends on the application.

However this doesn't mean that **Min-Max Scaling** is not useful at all. A popular application is **image processing**, where pixel intensities have to be normalized to fit within a certain range (i.e., $[0, 255]$ for the RGB colour range). Also, typical **Neural Network** Algorithm require data that are on a scale of 0 to 1.

a 0

⌄ Need for Normalization

For example, consider a data set containing two features, age(x1), and income(x2). Where age ranges from 0–100, while income ranges from 0–20,000 and higher. Income is about 1,000 times larger than age and ranges from 20,000–500,000. So, these two features are in very different ranges. When we do further analysis, like multivariate linear regression, for example, the attributed income will intrinsically influence the result more due to its larger value. But this doesn't necessarily mean it is more important as a predictor.

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

⌄ Building the ANN Model

```
# sequential model to initialise our ann and dense module to build the layers
from keras.models import Sequential
from keras.layers import Dense

classifier = Sequential()
# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu', input_dim = 10))

# Adding the second hidden layer
classifier.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))

# Adding the output layer
classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
```

⌄ Compiling and Fitting the Model

```
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])

# Fitting the ANN to the Training set
classifier.fit(X_train, y_train, batch_size = 10, epochs = 100, verbose = 1)
```

```
Epoch 78/100
700/700 [=====] - 1s 805us/step - loss: 0.3867 - accuracy: 0.8400
Epoch 79/100
700/700 [=====] - 1s 805us/step - loss: 0.3865 - accuracy: 0.8401
Epoch 80/100
700/700 [=====] - 1s 872us/step - loss: 0.3868 - accuracy: 0.8411
Epoch 81/100
700/700 [=====] - 1s 827us/step - loss: 0.3860 - accuracy: 0.8404
Epoch 82/100
700/700 [=====] - 1s 827us/step - loss: 0.3857 - accuracy: 0.8384
Epoch 83/100
700/700 [=====] - 1s 805us/step - loss: 0.3855 - accuracy: 0.8417
Epoch 84/100
700/700 [=====] - 1s 760us/step - loss: 0.3858 - accuracy: 0.8407
Epoch 85/100
700/700 [=====] - 1s 782us/step - loss: 0.3859 - accuracy: 0.8424
Epoch 86/100
700/700 [=====] - 1s 782us/step - loss: 0.3849 - accuracy: 0.8404
Epoch 87/100
700/700 [=====] - 1s 805us/step - loss: 0.3839 - accuracy: 0.8413
Epoch 88/100
700/700 [=====] - 1s 760us/step - loss: 0.3846 - accuracy: 0.8400
Epoch 89/100
```

▼ Testing the Model

```
score, acc = classifier.evaluate(X_train, y_train,
                                batch_size=10)
print('Train score:', score)
print('Train accuracy:', acc)

# Predicting the Test set results
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5)

print('*'*20)
score, acc = classifier.evaluate(X_test, y_test,
                                batch_size=10)
print('Test score:', score)
print('Test accuracy:', acc)

700/700 [=====] - 1s 782us/step - loss: 0.3342 - accuracy: 0.8631
Train score: 0.33418112993240356
Train accuracy: 0.8631428480148315
*****
300/300 [=====] - 0s 784us/step - loss: 0.3537 - accuracy: 0.8550
Test score: 0.35368815064430237
Test accuracy: 0.8550000190734863
```

▼ Confusion Matrix

* Accuracy

number of examples correctly predicted / total number of examples

$$ACC = (TP + TN) / (TP + FP + FN + TN)$$

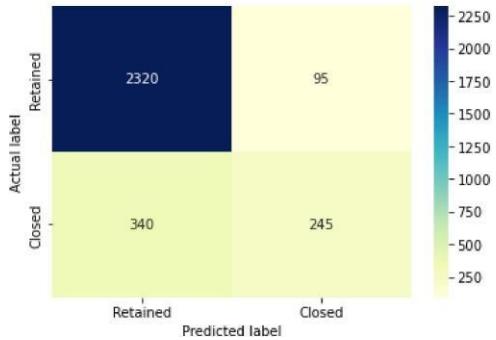
```
# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
target_names = ['Retained', 'Closed']
cm = confusion_matrix(y_test, y_pred)
print(cm)

[[2320  95]
 [ 340 245]]
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
p = sns.heatmap(pd.DataFrame(cm), annot=True, xticklabels=target_names, yticklabels=target_names, cmap="YlGnBu", fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')

Text(0.5, 15.0, 'Predicted label')
Confusion matrix
```



Classification Report

* True Positive Rate

number of samples actually and predicted as Positive / total number of samples actually Positive

Also called **Sensitivity or Recall**.

$$TPR = TP/P = TP/(TP+FN)$$

* Positive Predictive Value

number of samples actually and predicted as Positive / total number of samples predicted as Positive

Also called **Precision**.

$$PPV = TP/(TP+FP)$$

* F1 score

Harmonic Mean of Precision and Recall.

$$F_1 = \frac{2}{\frac{1}{precision} + \frac{1}{recall}} = 2 \cdot \frac{precision \cdot recall}{precision + recall}$$

```
#import classification_report
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred, target_names=target_names))

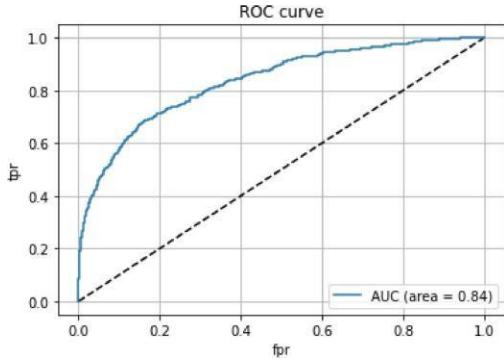
precision    recall   f1-score   support
Retained      0.87     0.96     0.91     2415
Closed         0.72     0.42     0.53      585
accuracy          --       --       --      3000
macro avg       0.80     0.69     0.72     3000
weighted avg    0.84     0.85     0.84     3000
```

ROC curve

```

from sklearn.metrics import roc_curve, auc
y_pred_proba = classifier.predict(X_test)
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
plt.plot([0,1],[0,1],'k--')
plt.plot(fpr,tpr, label='AUC (area = %0.2f)' % roc_auc)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.grid()
plt.legend(loc="lower right")
plt.title('ROC curve')
plt.show()

```



```

#Area under ROC curve
from sklearn.metrics import roc_auc_score
roc_auc_score(y_test,y_pred_proba)

0.8360007786094743

```

✓ Finetuning the Network

```

# Tuning the ANN
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
from keras.models import Sequential
from keras.layers import Dense
def build_classifier(optimizer):
    classifier = Sequential()
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu', input_dim = 10))
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu'))
    classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
    classifier.compile(optimizer = optimizer, loss = 'binary_crossentropy', metrics = ['accuracy'])
    return classifier
classifier = KerasClassifier(build_fn = build_classifier)
parameters = {'batch_size': [16, 32],
              'epochs': [50, 100],
              'optimizer': ['adam', 'rmsprop']}
grid_search = GridSearchCV(estimator = classifier,
                           param_grid = parameters,
                           scoring = 'accuracy', cv = 2) #,cv = 10
grid_search = grid_search.fit(X_train, y_train, verbose = 1)
best_parameters = grid_search.best_params_
best_accuracy = grid_search.best_score_

```



3/17/24, 7:39 PM

Exercise 1.ipynb - Colaboratory

```
Epoch 26/100
110/110 [=====] - 0s 1ms/step - loss: 0.4165 - accuracy: 0.8311
Epoch 27/100
110/110 [=====] - 0s 1ms/step - loss: 0.4159 - accuracy: 0.8320
Epoch 28/100
110/110 [=====] - 0s 1ms/step - loss: 0.4153 - accuracy: 0.8326
Epoch 29/100
110/110 [=====] - 0s 1ms/step - loss: 0.4146 - accuracy: 0.8326
Epoch 30/100
110/110 [=====] - 0s 1ms/step - loss: 0.4139 - accuracy: 0.8326
Epoch 31/100
110/110 [=====] - 0s 1ms/step - loss: 0.4134 - accuracy: 0.8349
Epoch 32/100
110/110 [=====] - 0s 1ms/step - loss: 0.4128 - accuracy: 0.8351
Epoch 33/100
110/110 [=====] - 0s 1ms/step - loss: 0.4123 - accuracy: 0.8360
Epoch 34/100
110/110 [=====] - 0s 1ms/step - loss: 0.4119 - accuracy: 0.8340
Epoch 35/100
110/110 [=====] - 0s 1ms/step - loss: 0.4115 - accuracy: 0.8360
Epoch 36/100
110/110 [=====] - 0s 1ms/step - loss: 0.4109 - accuracy: 0.8357
Epoch 37/100
110/110 [=====] - 0s 1ms/step - loss: 0.4106 - accuracy: 0.8380
Epoch 38/100
110/110 [=====] - 0s 1ms/step - loss: 0.4103 - accuracy: 0.8366
Epoch 39/100
110/110 [=====] - 0s 1ms/step - loss: 0.4098 - accuracy: 0.8371
Epoch 40/100
110/110 [=====] - 0s 1ms/step - loss: 0.4094 - accuracy: 0.8360
Epoch 41/100
110/110 [=====] - 0s 1ms/step - loss: 0.4090 - accuracy: 0.8380
Epoch 42/100
110/110 [=====] - 0s 1ms/step - loss: 0.4086 - accuracy: 0.8380
Epoch 43/100
110/110 [=====] - 0s 1ms/step - loss: 0.4083 - accuracy: 0.8377
Epoch 44/100
110/110 [=====] - 0s 1ms/step - loss: 0.4081 - accuracy: 0.8377
Epoch 45/100
110/110 [=====] - 0s 1ms/step - loss: 0.4078 - accuracy: 0.8374
Epoch 46/100
110/110 [=====] - 0s 1ms/step - loss: 0.4075 - accuracy: 0.8389
```

best_parameters

```
{'batch_size': 16, 'epochs': 50, 'optimizer': 'rmsprop'}
```

best_accuracy

Experiment : 2

AIM: Build an Artificial Neural Network to implement Binary Classification task using the Back-propagation algorithm and test the same using appropriate data sets.

Importing Necessary Libraries

```
In [ ]: import numpy as np  
import pandas as pd
```

Loading the Churn Dataset

This data set contains details of a bank's customers and the target variable is a binary variable reflecting the fact whether the customer left the bank (closed his account) or he continues to be a customer.

Binary flag 1 if the customer closed account with bank and 0 if the customer is retained.

```
In [ ]: churn_data = pd.read_csv('Churn_Modelling.csv', delimiter = ',', )  
churn_data.head(5)
```

```
Out[ ]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Ten
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

Accessing the Column Names in the Dataset

```
In [ ]: churn_data.columns
```

```
Out[ ]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
   'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
   'IsActiveMember', 'EstimatedSalary', 'Exited'],
  dtype='object')
```

Setting Column as a Index

```
In [ ]: churn_data = churn_data.set_index('RowNumber')
churn_data.head()
```

Finding the Shape of the Dataset

```
In [ ]: churn_data.shape
```

```
Out[ ]: (10000, 13)
```

```
In [ ]: churn_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 10000 entries, 1 to 10000
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
 ----  -- 
 0   CustomerId      10000 non-null   int64  
 1   Surname          10000 non-null   object  
 2   CreditScore      10000 non-null   int64  
 3   Geography         10000 non-null   object  
 4   Gender            10000 non-null   object  
 5   Age               10000 non-null   int64  
 6   Tenure            10000 non-null   int64  
 7   Balance           10000 non-null   float64 
 8   NumOfProducts     10000 non-null   int64  
 9   HasCrCard        10000 non-null   int64  
 10  IsActiveMember    10000 non-null   int64  
 11  EstimatedSalary   10000 non-null   float64 
 12  Exited            10000 non-null   int64  
dtypes: float64(2), int64(8), object(3)
memory usage: 1.1+ MB
```

Checking Missing Values

```
In [ ]: churn_data.isna().sum()
```

```
Out[ ]: CustomerId      0  
         Surname        0  
         CreditScore    0  
         Geography      0  
         Gender         0  
         Age            0  
         Tenure          0  
         Balance         0  
         NumOfProducts   0  
         HasCrCard       0  
         IsActiveMember  0  
         EstimatedSalary 0  
         Exited          0  
         dtype: int64
```

Some Columns are Totally Unproductive so let's Remove them

```
In [ ]: churn_data.nunique()
```

```
Out[ ]: CustomerId      10000  
         Surname        2932  
         CreditScore    460  
         Geography      3  
         Gender         2  
         Age            70  
         Tenure          11  
         Balance         6382  
         NumOfProducts   4  
         HasCrCard       2  
         IsActiveMember  2  
         EstimatedSalary 9999  
         Exited          2  
         dtype: int64
```

```
In [ ]: churn_data.drop(['CustomerId', 'Surname'], axis=1, inplace=True)
```

```
In [ ]: churn_data.head()
```

Out[]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfPro
RowNumber							
1	619	France	Female	42	2	0.00	
2	608	Spain	Female	41	1	83807.86	
3	502	France	Female	42	8	159660.80	
4	699	France	Female	39	1	0.00	
5	850	Spain	Female	43	2	125510.82	

In []: churn_data.shape

Out[]: (10000, 11)

Some Visualizations

```
In [ ]: from matplotlib import pyplot as plt
import seaborn as sns
from scipy import stats
df = churn_data.copy()
```

```
In [ ]: def plot_univariate(col):
    if(df[col].nunique() > 2):
        plt.figure(figsize=(10, 7))
        h = 0.15
        rot=90
    else:
        plt.figure(figsize=(6, 6))
        h = 0.5
        rot=0
    plot = sns.countplot(x = df[col], palette='pastel')

    for bars in plot.containers:
        for p in bars:
            plot.annotate(format(p.get_height()), (p.get_x() + p.get_width()/2, p.get_y() + h), ha = 'center', va = 'bottom')
            plot.annotate(f'{p.get_height() * 100 / df[col].shape[0]: .1f}%', (p.get_x() + p.get_width()/2, p.get_y() + h - 5), ha = 'center', va = 'bottom', rotation=rot)
```

```
In [ ]: def plot_bivariate(col, hue):
    if(df[col].nunique() > 5):
        plt.figure(figsize=(20, 10))
        rot=90
    else:
        plt.figure(figsize=(10, 7))
        rot=0
    def percentage(ax):
```

```

heights = [[p.get_height() for p in bars] for bars in ax.containers]
for bars in ax.containers:
    for i, p in enumerate(bars):
        total = sum(group[i] for group in heights) #Sum total of
        percentage = (100 * p.get_height() / total) #Calculate %
        ax.annotate(format(p.get_height()), (p.get_x() + p.get_wi
            ha = 'center', va = 'bottom', rotation=0)
        if(percentage>25.0):
            percentage = f'{percentage:.1f}%'#
            ax.annotate(percentage, (p.get_x() + p.get_width()*0
plot = sns.countplot(x=df[col], hue=df[hue], palette='pastel')
percentage(plot)

```

```

In [ ]: def spearman(df,hue):
    feature = []
    correlation = []
    result = []
    for col in df.columns:
        corr, p = stats.spearmanr(df[col], df[hue])
        feature.append(col)
        correlation.append(corr)
        alpha = 0.05
        if p > alpha:
            result.append('No correlation (fail to reject H0)')
        else:
            result.append('Some correlation (reject H0)')
    c = pd.DataFrame({'Feature Name':feature,'correlation coefficient':correlation})
    display(c)

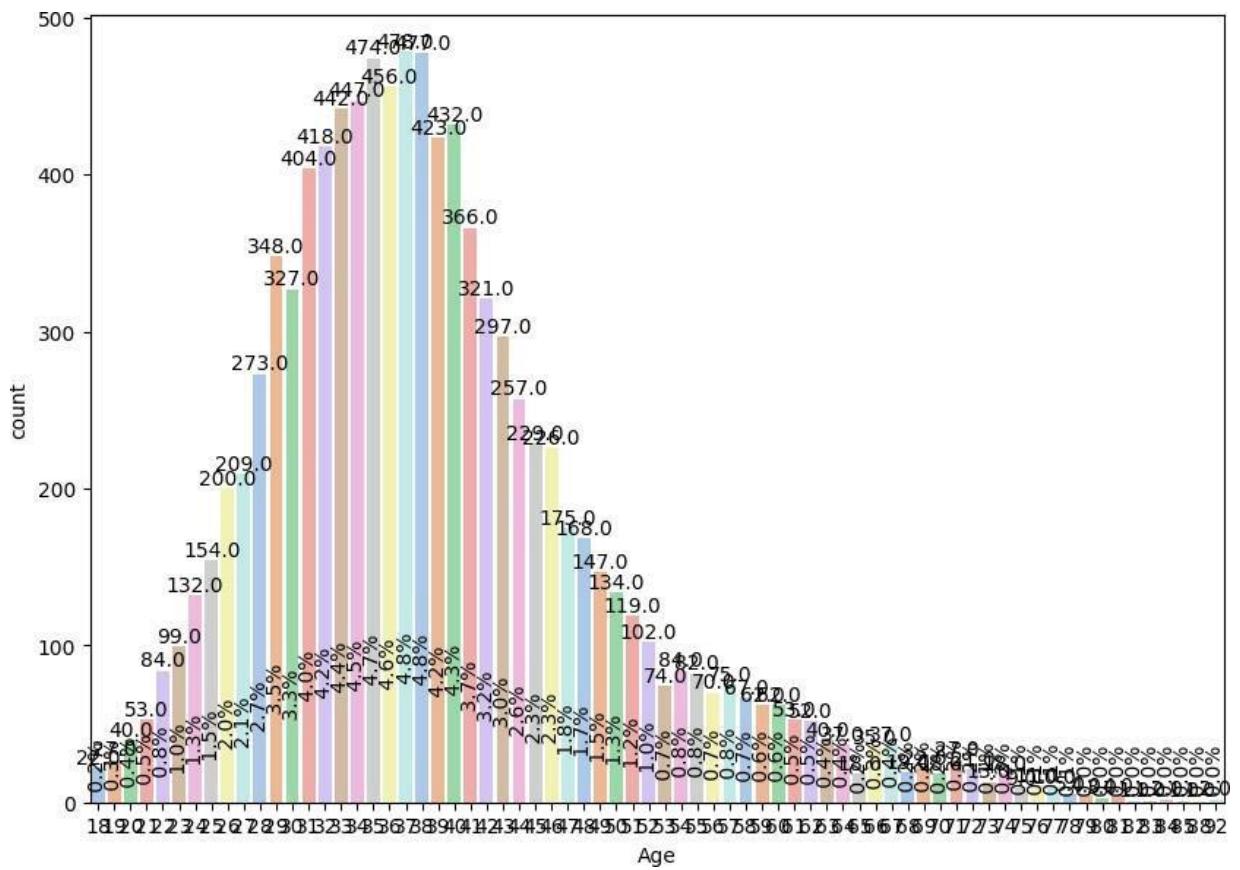
```

```
In [ ]: plot_univariate('Age')
```

```
/var/folders/sd/gsmj_sw13wj59mxmsl69t0bw0000gn/T/ipykernel_12551/420076724
8.py:10: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
plot = sns.countplot(x = df[col], palette='pastel')
```



```
In [ ]: plot_bivariate('Age', 'Exited')
```

```
In [ ]: spearman(churn_data, 'Age')
```

	Feature Name	correlation coefficient	Inference
0	CreditScore	-0.007974	No correlation (fail to reject H0)
1	Geography	0.035351	Some correlation (reject H0)
2	Gender	-0.029785	Some correlation (reject H0)
3	Age	1.000000	Some correlation (reject H0)
4	Tenure	-0.010405	No correlation (fail to reject H0)
5	Balance	0.033304	Some correlation (reject H0)
6	NumOfProducts	-0.058566	Some correlation (reject H0)
7	HasCrCard	-0.015278	No correlation (fail to reject H0)
8	IsActiveMember	0.039839	Some correlation (reject H0)
9	EstimatedSalary	-0.002431	No correlation (fail to reject H0)
10	Exited	0.323968	Some correlation (reject H0)

Label Encoding of Categorical Variables

Label Encoding means converting categorical features into numerical values. So that they can be fitted by machine learning models which only take numerical data.

Example: Suppose we have a column Height in some dataset that has elements as Tall, Medium, and short. To convert this categorical column into a numerical column we will apply label encoding to this column. After applying label encoding, the Height column is converted into a numerical column having elements 0,1, and 2 where 0 is the label for tall, 1 is the label for medium, and 2 is the label for short height.

```
In [ ]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
churn_data[['Geography', 'Gender']] = churn_data[['Geography', 'Gender']]
```

```
In [ ]: churn_data.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfPro
RowNumber							
1	619	0	0	42	2	0.00	
2	608	2	0	41	1	83807.86	
3	502	0	0	42	8	159660.80	
4	699	0	0	39	1	0.00	
5	850	2	0	43	2	125510.82	

Separating Label from Data

```
In [ ]: y = churn_data.Exited
x = churn_data.drop(['Exited'], axis=1)
```

```
In [ ]: x.columns
```

```
Out[ ]: Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
       'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary'],
       dtype='object')
```

```
In [ ]: y
```

Splitting the Data into Training and

Testing

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3)
```

```
In [ ]: print("Shape of the X_train", X_train.shape)
print("Shape of the X_test", X_test.shape)
print("Shape of the y_train", y_train.shape)
print("Shape of the y_test", y_test.shape)
```

Shape of the X_train (7000, 10)

Shape of the X_test (3000, 10)

Shape of the y_train (7000,)

Shape of the y_test (3000,)

Feature Scaling:

The result of standardization (or Z-Score normalization) is that the features will be re scaled so that they'll have the properties of a standard normal distribution with:

$$\mu = 0$$

And

$$\sigma = 1$$

Where μ is the mean(average) and σ is the standard deviation from the mean; standard scores (also called Z scores) of the samples are calculated as follows:

$$z = \frac{x - \mu}{\sigma}$$

About Min-Max Scaling

An alternative approach to Z-Score normalization (or called standardization) is the so-called Min-Max Scaling (often also simply called Normalization - a common cause for ambiguities)

In this approach, the data is scaled to a fixed range - usually $[0, 1]$. The cost of having this bounded range - in contrast to standardization - is that we will end up with smaller standard deviations, which can suppress the effect of outliers.

Note:

If the dataset has lots of outliers, and the algorithms are sensitive to outliers,

please use Min-Max Scaler

A Min-Max Scaling is typically done via the following equation:

$$X_{norm} = \frac{X_i - X_{min}}{X_{max} - X_{min}}$$

X_i is the i^{th} sample of dataset.

Z-Score Standardization or Min-Max Scaling

"Standardization or Min-Max scaling"? - There is no obvious answer to this question: it really depends on the application.

However this doesn't mean that Min-Max Scaling is not useful at all, A popular application is image processing , where pixel intensities have to be normalized to fit within a certain range (i.e., [0, 255] for the RGB colour range). Also, typical Neural Network Algorithm require data that on a 0 – 1 scale.

Need for Normalization

For example, consider a data set containing two features, age(x_1), and income(x_2). Where age ranges from 0-100, while income ranges from 0-20,000 and higher. Income is about 1,000 times larger than age and ranges from 20,000-500,000. So, these two features are in very different ranges. When we do further analysis, like multivariate linear regression, for example, the attributed income will intrinsically influence the result more due to its larger value. But this doesn't necessarily mean it is more important as a predictor.

```
In [ ]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

Building the ANN Model

```
In [ ]: # sequential model to initialise our ann and dense module to build the la
from keras.models import Sequential
from keras.layers import Dense
```

```
In [ ]: classifier = Sequential()
# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation='relu'))

# Adding the second hidden layer
classifier.add(Dense(units = 16, kernel_initializer = 'uniform', activation='relu'))

# Adding the output layer
classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation='sigmoid'))
```

Compiling and Fitting the Model

```
In [ ]: classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics=['accuracy'])

# Fitting the ANN to the Training set
classifier.fit(X_train, y_train, batch_size = 10, epochs = 100, verbose = 1)
```

Testing the Model

```
In [ ]: score, acc = classifier.evaluate(X_train, y_train,
                                         batch_size=10)
print('Train score:', score)
print('Train accuracy:', acc)

# Predicting the Test set results
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5)

print('*'*20)
score, acc = classifier.evaluate(X_test, y_test,
                                         batch_size=10)
print('Test score:', score)
print('Test accuracy:', acc)
```

```
700/700 [=====] - 0s 278us/step - loss: 0.3287 - accuracy: 0.8637
Train score: 0.328686386346817
Train accuracy: 0.8637142777442932
94/94 [=====] - 0s 277us/step
*****
300/300 [=====] - 0s 300us/step - loss: 0.3470 - accuracy: 0.8593
Test score: 0.34699997305870056
Test accuracy: 0.859333336353302
```

Confusion Matrix

* Accuracy

number of examples correctly predicted / total number of examples

$$ACC = (TP + TN) / (TP + FP + FN + TN)$$

```
In [ ]: # Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
target_names = ['Retained', 'Closed']
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

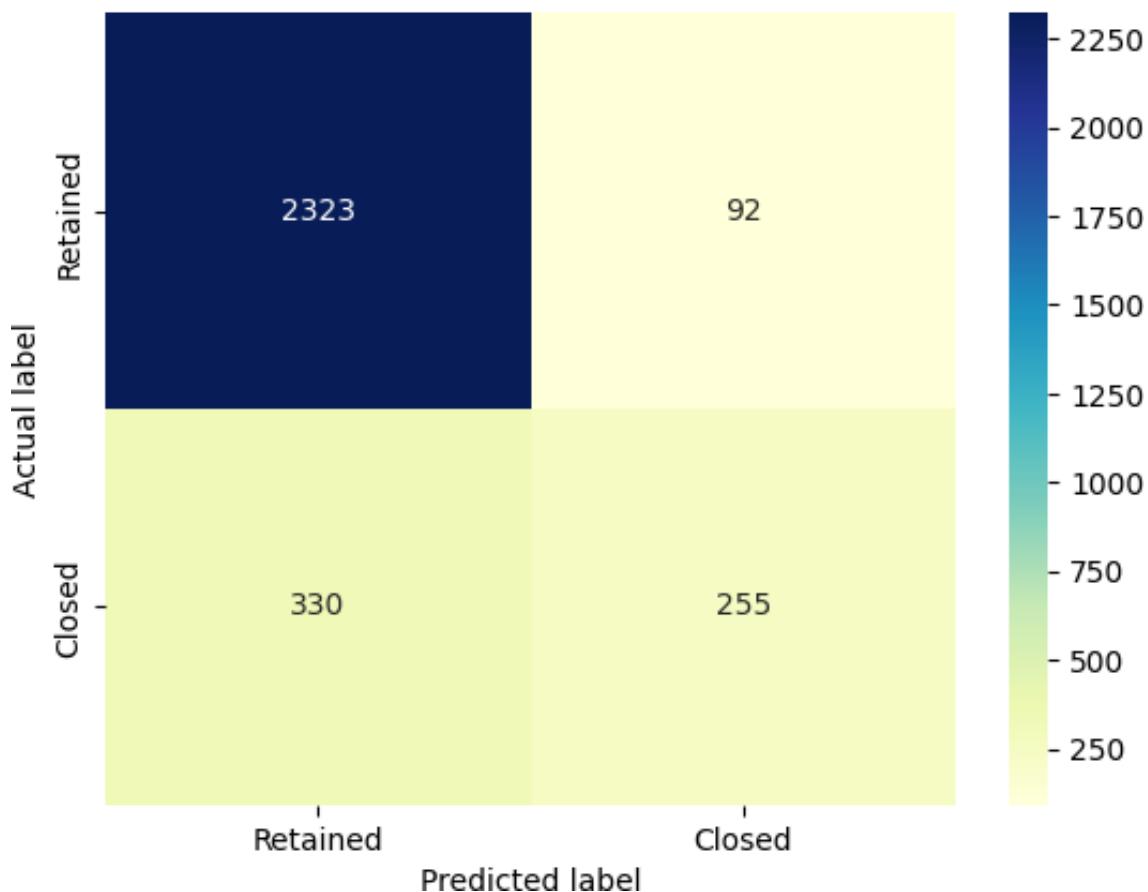
```
[2323    92]
[ 330   255]
```

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: p = sns.heatmap(pd.DataFrame(cm), annot=True, xticklabels=target_names, y
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
Out[ ]: Text(0.5, 23.522222222222222, 'Predicted label')
```

Confusion matrix



Classification Report

* True Positive Rate

number of samples actually and predicted as **Positive** / total number of samples actually **Positive**

Also called Sensitivity or Recall.

$$TPR = TP/P = TP/(TP + FN)$$

* Positive Predictive Value

number of samples actually and predicted as **Positive** / total number of samples predicted as **Positive**

Also called Precision.

$$PPV = TP/(TP + FP)$$

* F1 score

Harmonic Mean of Precision and Recall.

$$F_1 = \frac{2}{\frac{1}{\text{recall}} + \frac{1}{\text{precision}}} = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

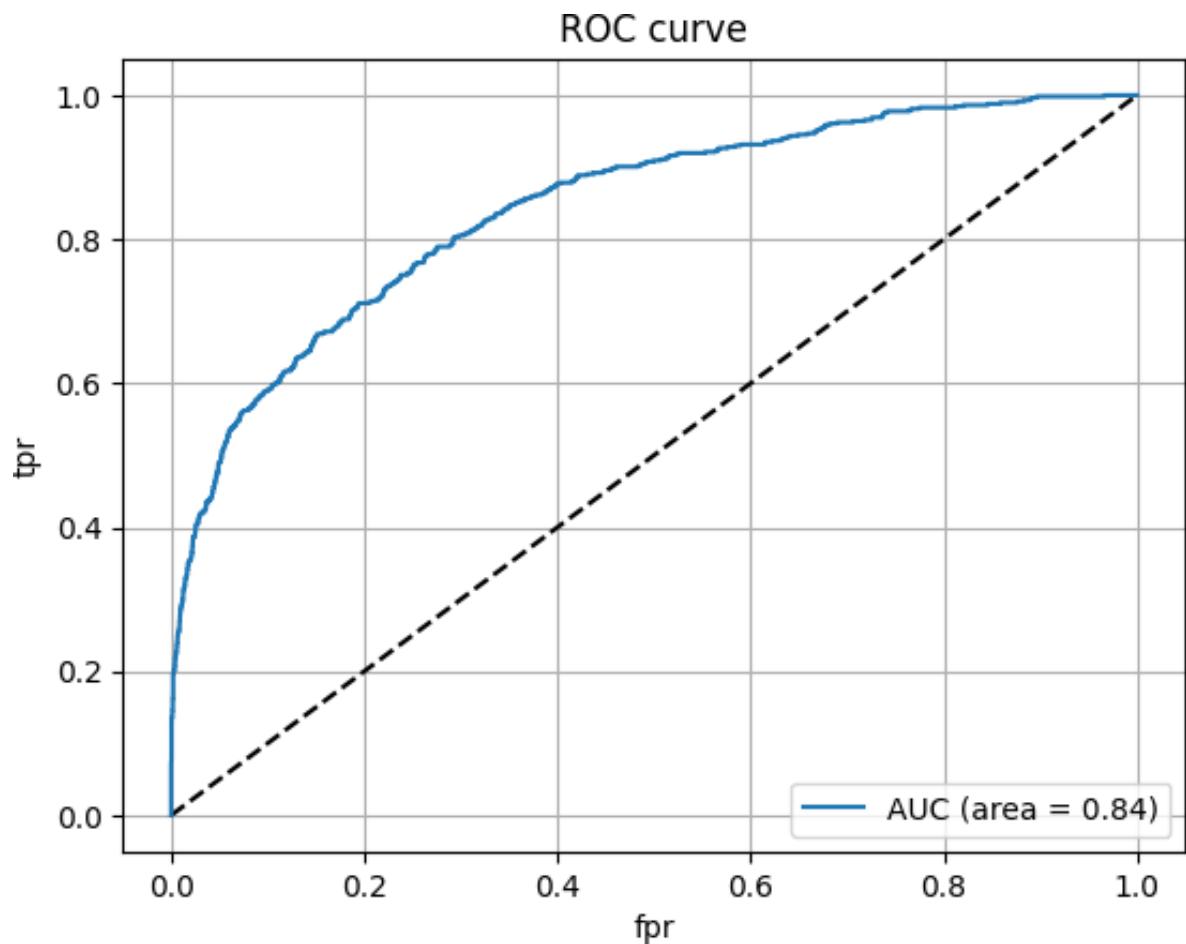
```
In [ ]: #import classification_report
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
Retained	0.88	0.96	0.92	2415
Closed	0.73	0.44	0.55	585
accuracy			0.86	3000
macro avg	0.81	0.70	0.73	3000
weighted avg	0.85	0.86	0.84	3000

ROC curve

```
In [ ]: from sklearn.metrics import roc_curve, auc
y_pred_proba = classifier.predict(X_test)
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
plt.plot([0,1],[0,1],'k--')
plt.plot(fpr,tpr, label='AUC (area = %0.2f)' % roc_auc)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.grid()
plt.legend(loc="lower right")
plt.title('ROC curve')
plt.show()
```

94/94 [=====] - 0s 314us/step



```
In [ ]: #Area under ROC curve
from sklearn.metrics import roc_auc_score
roc_auc_score(y_test,y_pred_proba)
```

Out[]: 0.8436201801419193

Experiment : 3

AIM: Build an Artificial Neural Network to implement Multi-Class Classification task using the Back-propagation algorithm and test the same using appropriate data sets.

Importing Necessary Libraries

```
In [ ]: import numpy as np  
import pandas as pd
```

Loading the Combined Cycle Power Plant Dataset

The dataset contains 9568 data points collected from a Combined Cycle Power Plant over 6 years (2006-2011), when the power plant was set to work with full load. Features consist of hourly average ambient variables Temperature (AT), Ambient Pressure (AP), Relative Humidity (RH) and Exhaust Vacuum (V) to predict the net hourly electrical energy output (PE) of the plant.

- For more information visit below link: Link:
<https://archive.ics.uci.edu/ml/datasets/Combined+Cycle+Power+Plant>

```
In [ ]: Powerplant_data = pd.read_excel('Folds5x2_pp.xlsx')  
Powerplant_data.head(5)
```

```
Out[ ]:      AT      V      AP      RH      PE  
0   14.96  41.76  1024.07  73.17  463.26  
1   25.18  62.96  1020.04  59.08  444.37  
2    5.11  39.40  1012.16  92.14  488.56  
3   20.86  57.32  1010.24  76.64  446.48  
4   10.82  37.50  1009.23  96.62  473.90
```

Accessing the Column Names in the

Dataset

```
In [ ]: Powerplant_data.columns
```

```
Out[ ]: Index(['AT', 'V', 'AP', 'RH', 'PE'], dtype='object')
```

Finding the Shape of the Dataset

```
In [ ]: Powerplant_data.shape
```

```
Out[ ]: (9568, 5)
```

```
In [ ]: Powerplant_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9568 entries, 0 to 9567
Data columns (total 5 columns):
 #   Column   Non-Null Count   Dtype  
 --- 
 0   AT        9568 non-null    float64
 1   V         9568 non-null    float64
 2   AP        9568 non-null    float64
 3   RH        9568 non-null    float64
 4   PE        9568 non-null    float64
dtypes: float64(5)
memory usage: 373.9 KB
```

Checking Missing Values

```
In [ ]: Powerplant_data.isna().sum()
```

```
Out[ ]: AT      0
        V       0
        AP     0
        RH     0
        PE     0
       dtype: int64
```

Number of Unique Items in Each Column

```
In [ ]: Powerplant_data.nunique()
```

```
Out[ ]: AT      2773
         V       634
         AP     2517
         RH     4546
         PE     4836
        dtype: int64
```

Seperating Label from Data

```
In [ ]: X = Powerplant_data.iloc[:, :-1].values
y = Powerplant_data.iloc[:, -1].values
```

or

```
In [ ]: # y o Powerplant_data['PE']
# X o Powerplant_data.drop(['PE'], axis=1)
```

Splitting the Data into Training and Testing

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_s
```

```
In [ ]: print("Shape of the X_train", X_train.shape)
print("Shape of the X_test", X_test.shape)
print("Shape of the X_val", X_val.shape)
print("Shape of the y_train", y_train.shape)
print("Shape of the y_test", y_test.shape)
print("Shape of the y_val", y_val.shape)
```

```
Shape of the X_train (6123, 4)
Shape of the X_test (1914, 4)
Shape of the X_val (1531, 4)
Shape of the y_train (6123,)
Shape of the y_test (1914,)
Shape of the y_val (1531,)
```

Building the ANN Model

```
In [ ]: # sequential model to initialise our ann and dense module to build the la
from keras.models import Sequential
from keras.layers import Dense
```

```
In [ ]: classifier = Sequential()
```

```
# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation='relu'))

# Adding the second hidden layer
classifier.add(Dense(units = 16, kernel_initializer = 'uniform', activation='relu'))

# Adding the third hidden layer
classifier.add(Dense(units = 32, kernel_initializer = 'uniform', activation='relu'))

# Adding the output layer
classifier.add(Dense(units = 1, kernel_initializer = 'uniform'))
```

Compiling and Fitting the Model

```
In [ ]: classifier.compile(optimizer = 'adam', loss = 'mean_squared_error', metrics=['mae'])

# Fitting the ANN to the Training set
model = classifier.fit(X_train, y_train, batch_size = 32, epochs = 200, validation_split=0.2,
                        shuffle=True)
```

Testing the Model

```
In [ ]: y_pred = classifier.predict(X_test)
np.set_printoptions(precision=2)
print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1)), axis=1))

60/60 [=====] - 0s 293us/step
[[463.95 471.81]
 [443.18 444.61]
 [437.15 437.03]
 ...
 [440.91 442.77]
 [441.47 445.33]
 [432.53 432.94]]
```

Metric values

MAE (Mean Absolute Error) :-

$$MAE = \left(\frac{1}{n} \right) \sum_{i=1}^n |y_i - \hat{y}_i|$$

MSE (Mean Square Error) :-

$$MSE = \left(\frac{1}{n} \right) \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

where y = actual value in the data set ; \hat{y} = value computed by solving the regression equation

RMSE (Root Mean Square Error) :-

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

```
In [ ]: import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [ ]: import sklearn.metrics
from math import sqrt
mae_no = sklearn.metrics.mean_absolute_error(y_test, classifier.predict(X_t)
mse_no = sklearn.metrics.mean_squared_error(y_test, classifier.predict(X_t)
rms = sqrt(sklearn.metrics.mean_squared_error(y_test, classifier.predict(X_t))

60/60 [=====] - 0s 678us/step
60/60 [=====] - 0s 263us/step
60/60 [=====] - 0s 269us/step
```

```
In [ ]: print('Mean Absolute Error      : ', mae_no)
print('Mean Square Error       : ', mse_no)
print('Root Mean Square Error: ', rms)
```

```
Mean Absolute Error      : 4.156925082630249
Mean Square Error       : 28.931698846583828
Root Mean Square Error: 5.37881946588504
```

RESULT: We have successfully built an Artificial Neural Network to implement Multi-Class Classification task using the Back-propagation algorithm and test the same using appropriate data sets.

Experiment : 4

AIM: Implement an image classification task using pre-trained models like AlexNet, VGGNet, InceptionNet and ResNet and compare the results.

Importing Necessary libraries

- From keras library we are going to use image preprocessing task, to normalize the image pixel values in between 0 to 1.
- Model is imported to load various Neural Network models such as Sequential.
- We are going to use transfer learning technique.

```
In [ ]: import keras
import numpy as np
from keras import Input
from keras import models
from keras import layers
from keras import optimizers
from keras.models import Model
from keras import applications
from keras import backend as k
import matplotlib.pyplot as plt
from keras.optimizers import SGD, Adam
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
from keras.preprocessing import image
from keras.models import Sequential, Model
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import Dropout, Flatten, Dense, GlobalAveragePooling2D
from keras.callbacks import ModelCheckpoint, LearningRateScheduler, TensorBoard
```

Loading the Training and Testing Data and Defining the Basic Parameters

- We are resizing the input image to 128 * 128
- In the dataset : Training Set : 70% Validation Set : 20% Test Set : 10%

```
In [ ]: # Normalize training and validation data in the range of 0 to 1
train_datagen = ImageDataGenerator(rescale=1./255) # vertical_flip=True,
# horizontal_flip=True
```

```
# height_shift_range=0
# width_shift_range=0.

validation_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

# Read the training sample and set the batch size
train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/plant_v1',
    target_size=(128, 128),
    batch_size=16,
    class_mode='categorical')

# Read Validation data from directory and define target size with batch size
validation_generator = validation_datagen.flow_from_directory(
    '/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/plant_v1',
    target_size=(128, 128),
    batch_size=16,
    class_mode='categorical',
    shuffle=False)

test_generator = test_datagen.flow_from_directory(
    '/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/plant_v1',
    target_size=(128, 128),
    batch_size=1,
    class_mode='categorical',
    shuffle=False)
```

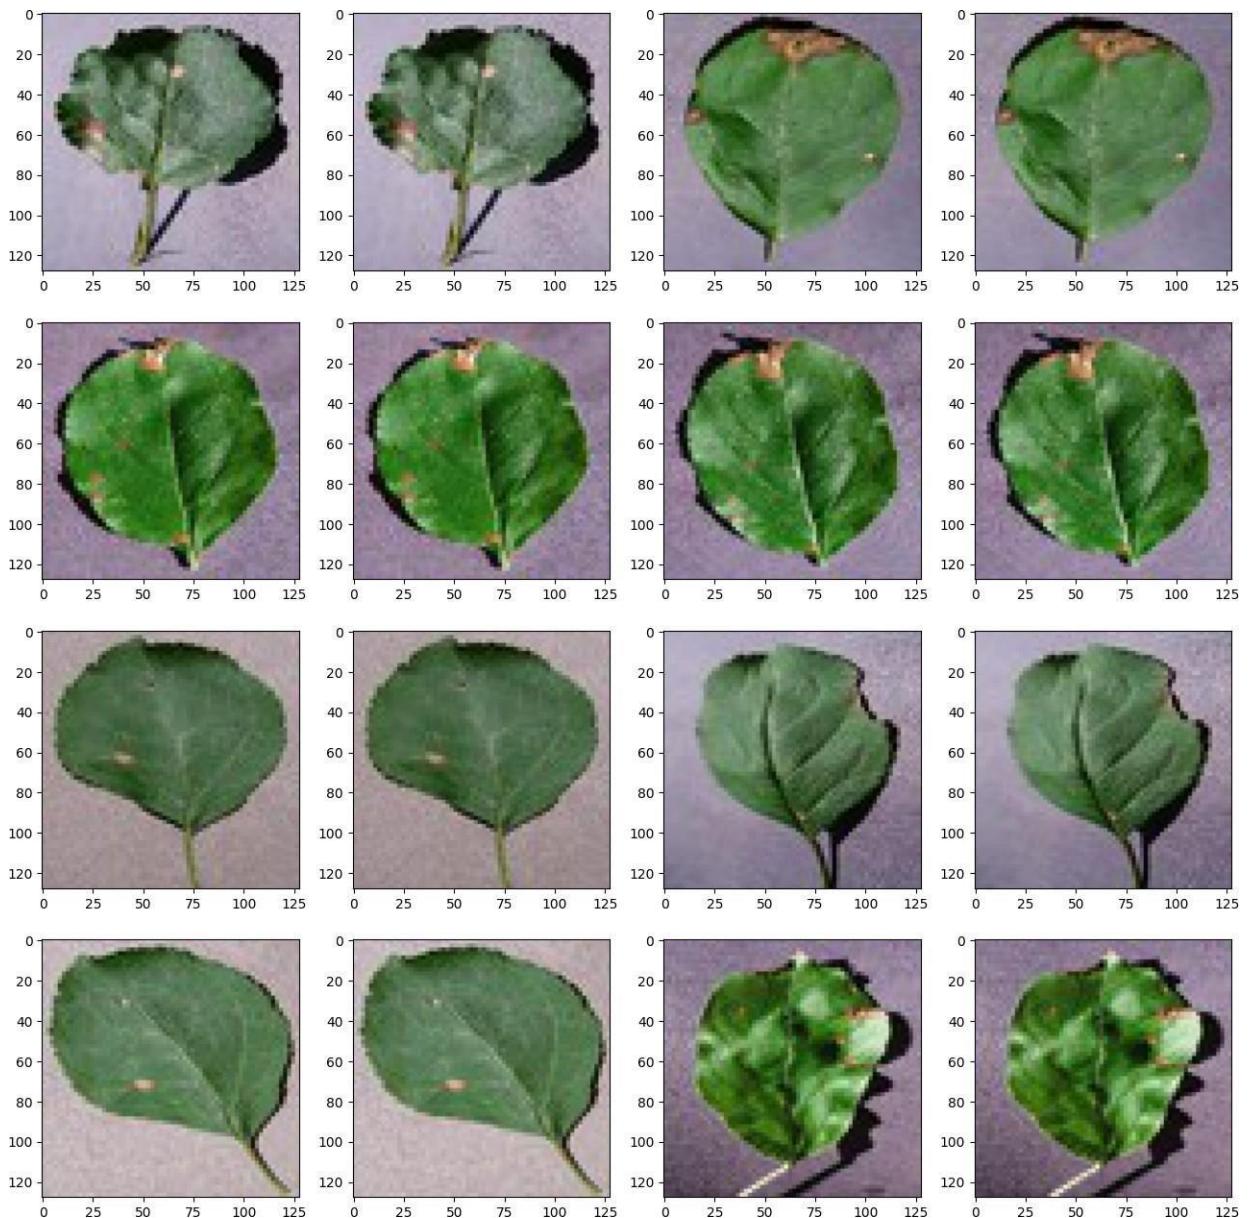
Found 3002 images belonging to 4 classes.

Found 635 images belonging to 4 classes.

Found 546 images belonging to 4 classes.

Visualization of Few Images

```
In [ ]: plt.figure(figsize=(16, 16))
for i in range(1, 17):
    plt.subplot(4, 4, i)
    img, label = test_generator.next()
    # print(img.shape)
    # print(label)
    plt.imshow(img[0])
plt.show()
```



```
In [ ]: img, label = test_generator.next()  
img[0].shape
```

```
Out[ ]: (128, 128, 3)
```

What is ImageNet?

ImageNet is formally a project aimed at (manually) labeling and categorizing images into almost 22,000 separate object categories for the purpose of computer vision research.

However, when we hear the term “ImageNet” in the context of deep learning and Convolutional Neural Networks, we are likely referring to the ImageNet Large Scale Visual Recognition Challenge, or ILSVRC for short.

The goal of this image classification challenge is to train a model that can correctly classify an input image into 1,000 separate object categories.

Models are trained on ~1.2 million training images with another 50,000 images for validation and 100,000 images for testing.

Exploring Keras Applications for Transfer Learning

VGG16

```
In [ ]: from tensorflow.keras.applications.vgg16 import VGG16

## Loading VGG16 model
base_model = VGG16(weights="imagenet", include_top=False, input_shape=(1
base_model.trainable = False ## Not trainable weights

base_model.summary()
```

Adding top layers according to number of classes in our data

```
In [ ]: flatten_layer = layers.GlobalAveragePooling2D()
# dense_layer_1 = layers.Dense(63, activation='relu')
# dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')

model = models.Sequential([
    base_model,
    flatten_layer,
    prediction_layer
])

model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
<hr/>		
vgg16 (Functional)	(None, 4, 4, 512)	14714688
global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 512)	0
dense_2 (Dense)	(None, 4)	2052
<hr/>		
Total params: 14,716,740		
Trainable params: 2,052		
Non-trainable params: 14,714,688		

Training

```
In [ ]: # sgd = SGD(lr=0.001, decay=1e-6, momentum=0.9, nesterov=True)
# We are going to use accuracy metrics and cross entropy loss as performance
model.compile(optimizer = Adam(learning_rate = 0.001), loss='categorical_
# Train the model
history = model.fit(train_generator,
    steps_per_epoch=train_generator.samples/train_generator.batch_size,
    epochs=30,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples/validation_generator.
    verbose=1)
```

Saving the model

```
In [ ]: model.save("VGG16_plant_deseas.h5")
print("Saved model to disk")
```

Saved model to disk

Loading the model

```
In [ ]: model = models.load_model('VGG16_plant_deseas.h5')
print("Model is loaded")
```

Model is loaded

Saving the Weights

```
In [ ]: model.save_weights('cnn_classification.h5')
```

Loading the weights

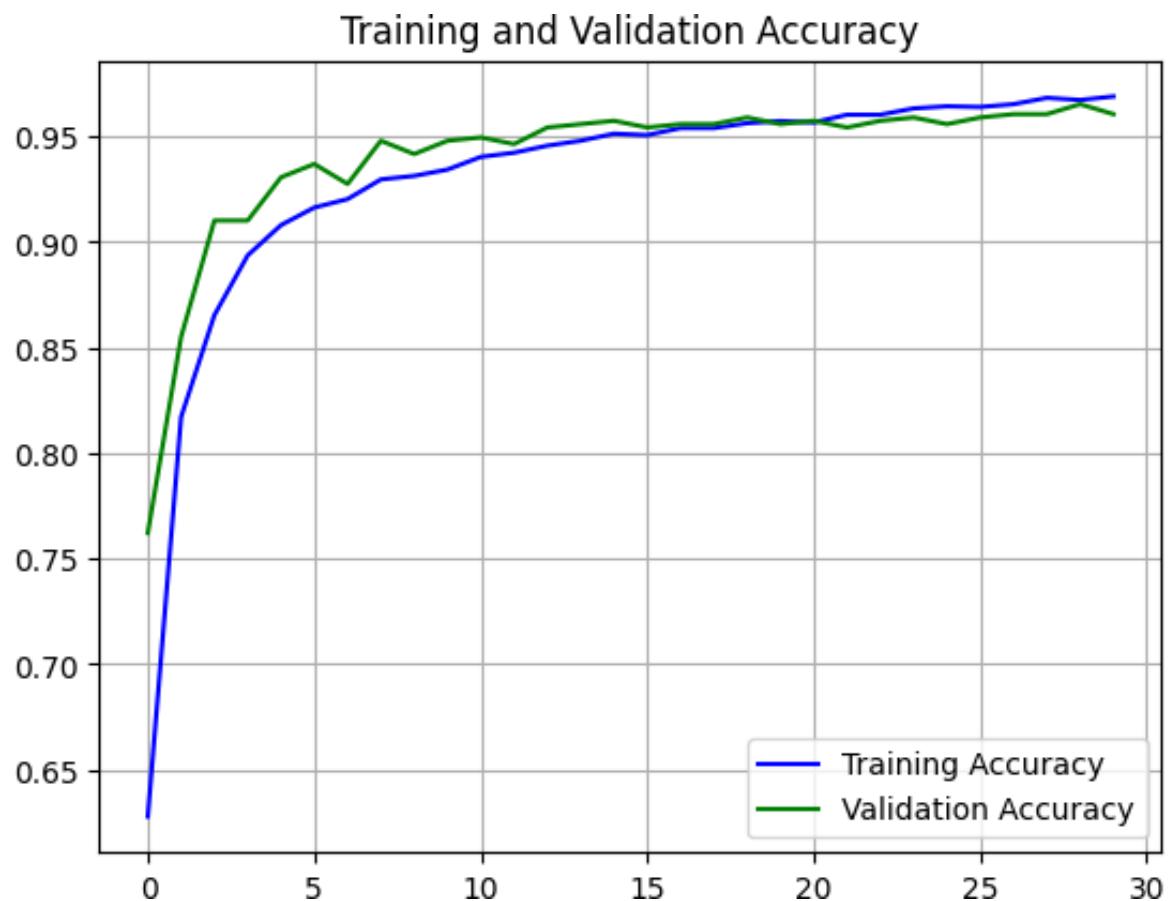
```
In [ ]: model.load_weights('cnn_classification.h5')
```

Visualization of training over epoch

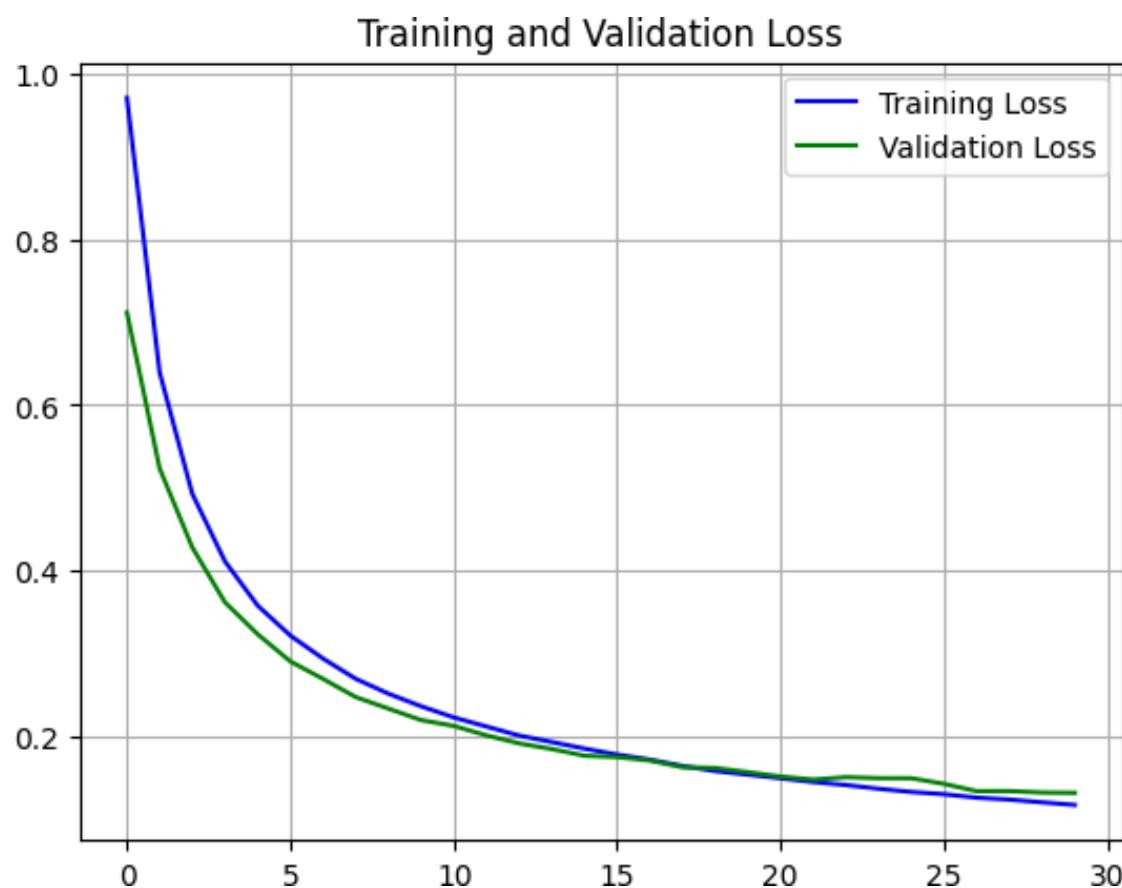
```
In [ ]: train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']
```

```
In [ ]: epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```



<Figure size 640x480 with 0 Axes>



Performance measure

```
In [ ]: # Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples)
predicted_classes = np.argmax(predictions, axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {} / {}".format(len(errors), test_generator.samples))
```

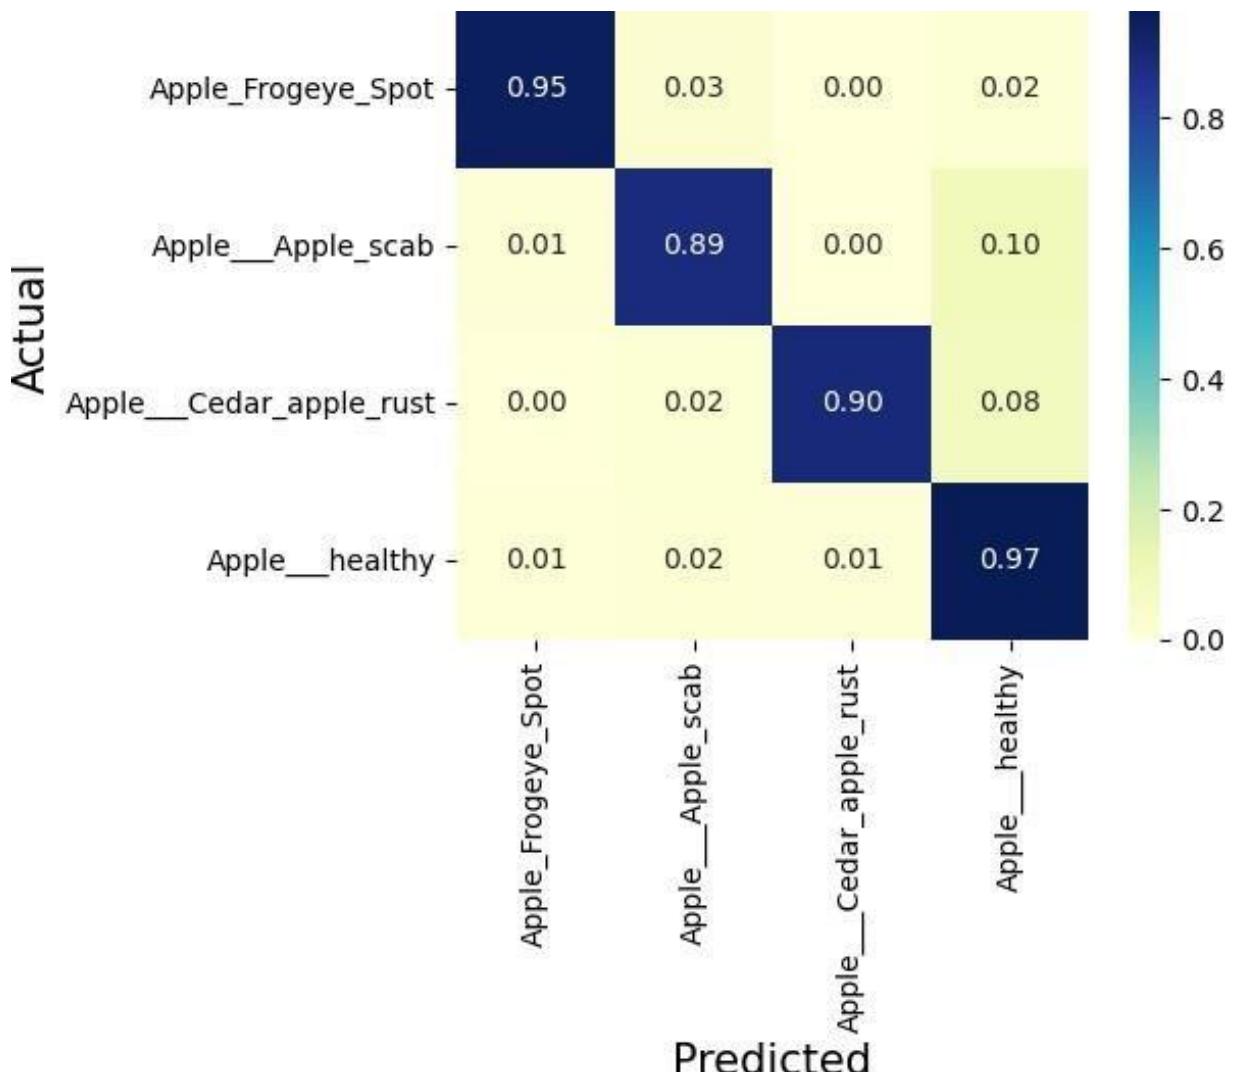
<ipython-input-20-1bdf1cd647ce>:14: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.

```
predictions = model.predict_generator(test_generator, steps=test_generator.samples/test_generator.batch_size, verbose=1)
546/546 [=====] - 3s 5ms/step
No of errors = 34/546
```

```
In [ ]: accuracy = ((test_generator.samples - len(errors)) / test_generator.samples)
accuracy
```

Out[]: 93.77289377289377

```
In [ ]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5, 4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index, yticklabels=label2index)
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)
```



```
In [ ]: from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names
```

	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.96	0.95	0.96	103
Apple_Apple_scab	0.94	0.89	0.91	134
Apple_Cedar_apple_rust	0.94	0.90	0.92	49
Apple_healthy	0.93	0.97	0.95	260
accuracy			0.94	546
macro avg	0.94	0.93	0.93	546
weighted avg	0.94	0.94	0.94	546

```
In [ ]:
```

InceptionNet

```
In [ ]: from keras import applications
```

```
## Loading InceptionV3 model
base_model = applications.InceptionV3(weights="imagenet", include_top=False)
base_model.trainable = False ## Not trainable weights

base_model.summary()
```

```
In [ ]: flatten_layer = layers.GlobalAveragePooling2D()
# dense_layer_1 = layers.Dense(63, activation='relu')
# dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')

model = models.Sequential([
    base_model,
    flatten_layer,
    # dense_layer_1,
    # dense_layer_2,
    prediction_layer
])

model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
<hr/>		
inception_v3 (Functional)	(None, 2, 2, 2048)	21802784
global_average_pooling2d_3 (GlobalAveragePooling2D)	(None, 2048)	0
dense_3 (Dense)	(None, 4)	8196
<hr/>		
Total params: 21,810,980		
Trainable params: 8,196		
Non-trainable params: 21,802,784		

```
In [ ]: model.compile(optimizer = Adam(learning_rate = 0.001), loss='categorical_crossentropy'
# Train the model
history = model.fit(train_generator,
                     steps_per_epoch=train_generator.samples/train_generator.batch_size,
                     epochs=30,
                     validation_data=validation_generator,
                     validation_steps=validation_generator.samples/validation_generator.batch_size,
                     verbose=1)
```

```
In [ ]: model.save("InceptionNet_plant_deseas.h5")
print("Saved model to disk")
```

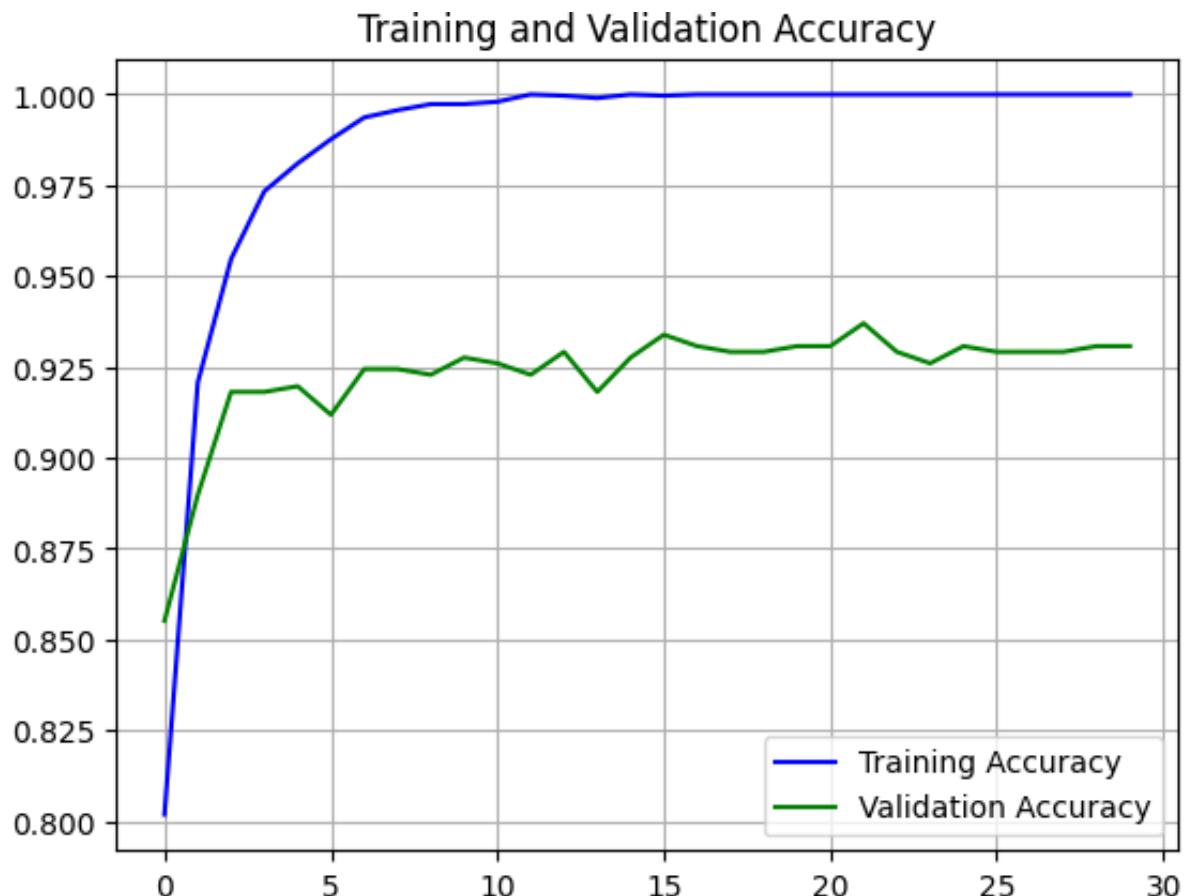
Saved model to disk

```
In [ ]: train_acc = history.history['acc']
```

```
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']
```

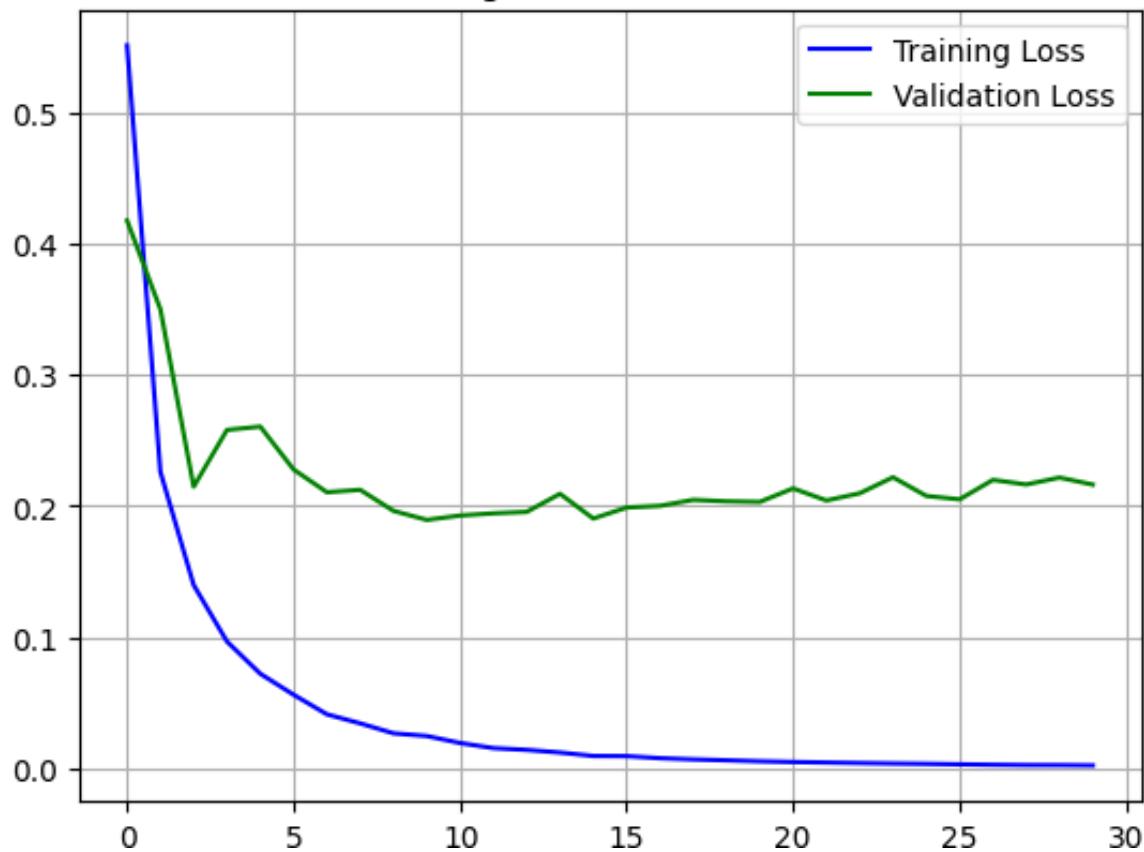
```
In [ ]: epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```



<Figure size 640x480 with 0 Axes>

Training and Validation Loss



```
In [ ]: # Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples)
predicted_classes = np.argmax(predictions, axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}/{}".format(len(errors), test_generator.samples))
```

<ipython-input-30-1bdf1cd647ce>:14: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.

```
predictions = model.predict_generator(test_generator, steps=test_generator.samples/test_generator.batch_size, verbose=1)
```

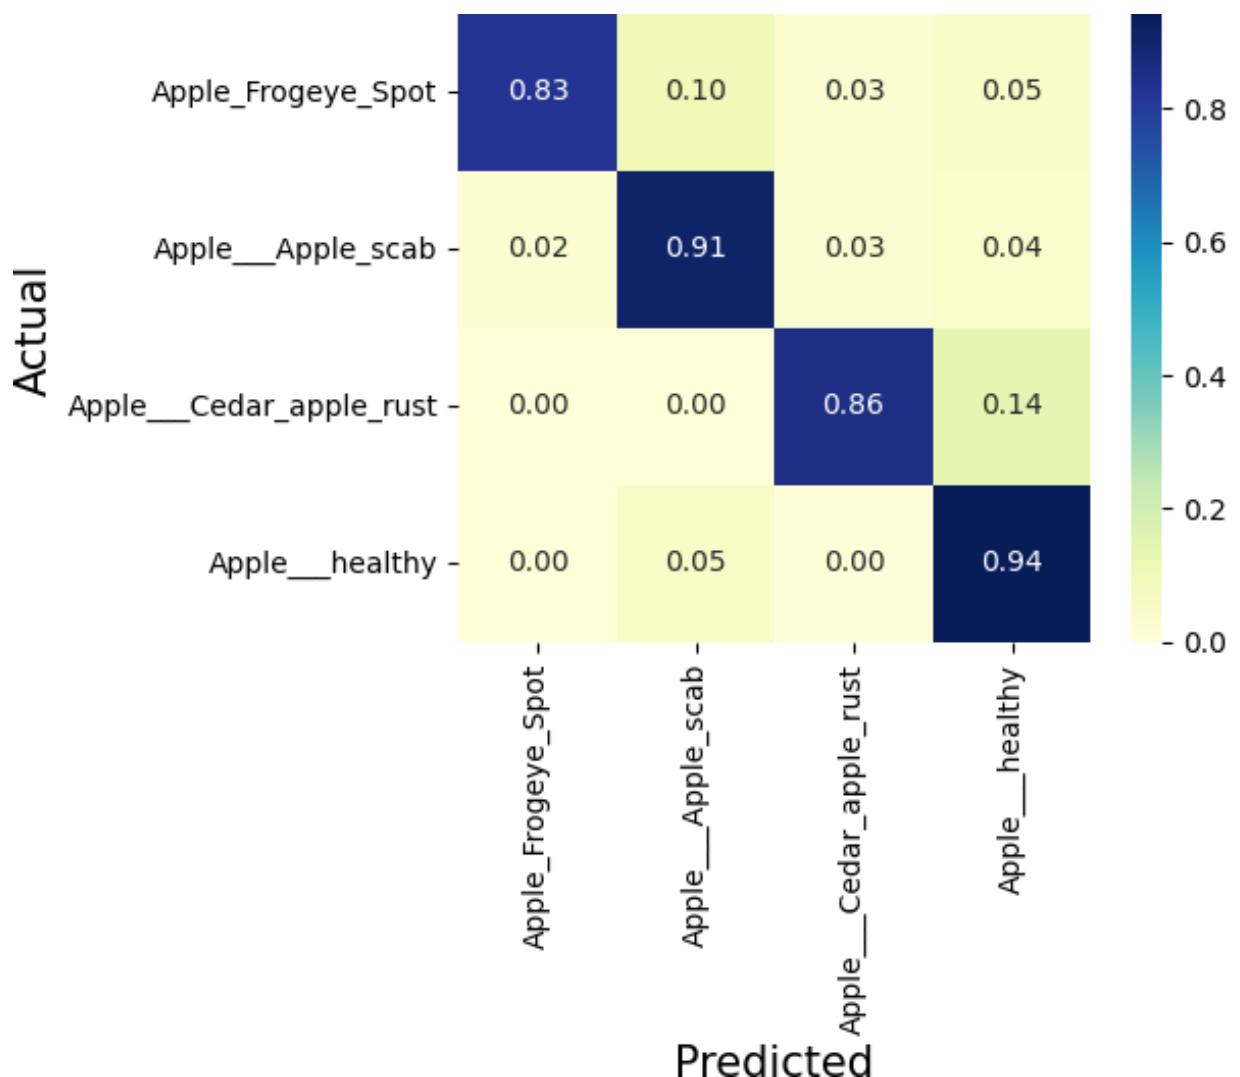
```
546/546 [=====] - 8s 11ms/step
```

```
No of errors = 52/546
```

```
In [ ]: accuracy = ((test_generator.samples - len(errors)) / test_generator.samples)
accuracy
```

Out[]: 90.47619047619048

```
In [ ]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5, 4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index, yticklabels=label2index,
            plt.ylabel('Actual', fontsize=15),
            plt.xlabel('Predicted', fontsize=15))
plt.show(block=False)
```



```
In [ ]: from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names
```

	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.97	0.83	0.89	103
Apple_Apple_scab	0.84	0.91	0.87	134
Apple_Cedar_apple_rust	0.84	0.86	0.85	49
Apple_healthy	0.94	0.94	0.94	260
accuracy			0.90	546
macro avg	0.89	0.88	0.89	546
weighted avg	0.91	0.90	0.90	546

In []:

ResNet

In []:

```
from keras import applications
```

```
## Loading VGG16 model
base_model = applications.ResNet50(weights="imagenet", include_top=False,
base_model.trainable = False ## Not trainable weights

base_model.summary()
```

In []:

```
flatten_layer = layers.GlobalAveragePooling2D()
# dense_layer_1 = layers.Dense(63, activation='relu')
# dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')
```

```
model = models.Sequential([
    base_model,
    flatten_layer,
    # dense_layer_1,
    # dense_layer_2,
    prediction_layer
])
```

```
model.summary()
```

In []:

```
model.compile(optimizer = Adam(learning_rate = 0.001), loss='categorical_
# Train the model
history = model.fit(train_generator,
    steps_per_epoch=train_generator.samples/train_generator.batch_size,
    epochs=30,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples/validation_generator.
    verbose=1)
```

In []:

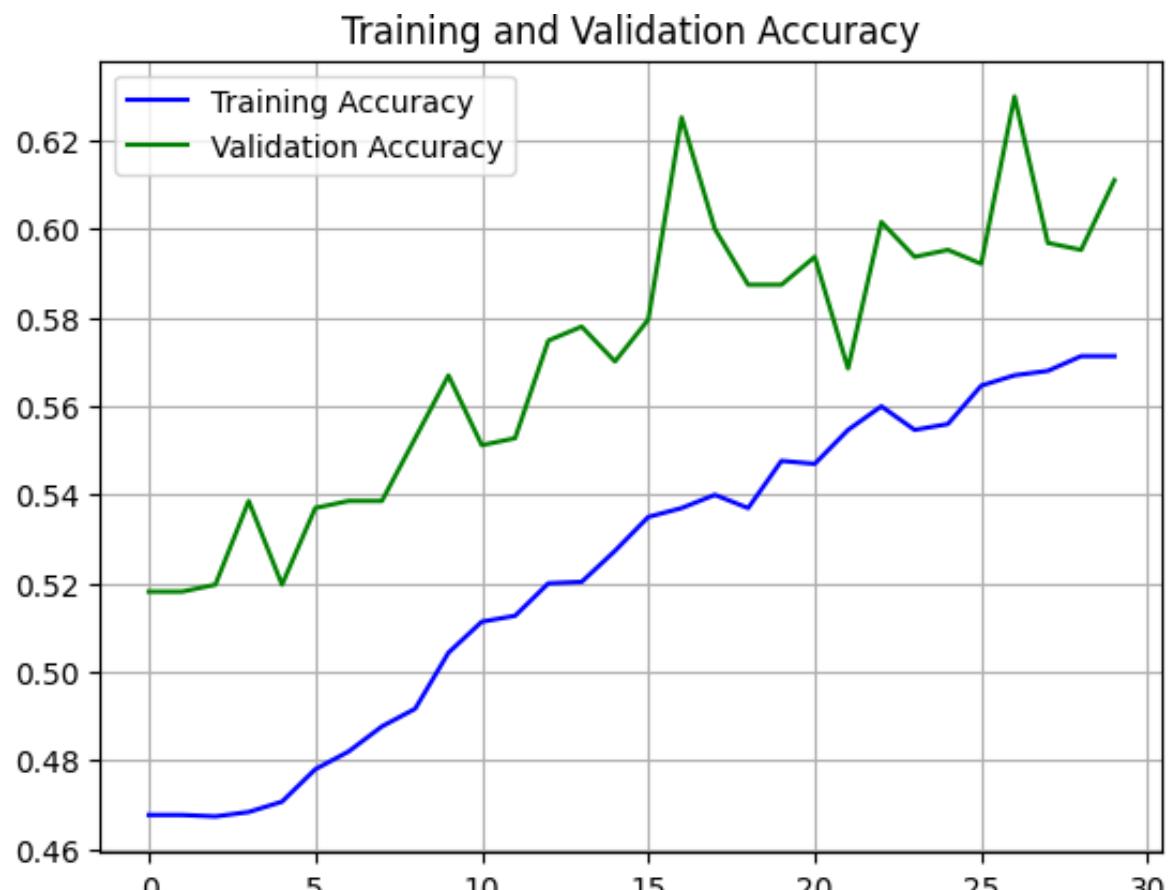
```
model.save("ResNet_plant_deseas.h5")
print("Saved model to disk")
```

Saved model to disk

```
In [ ]: train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']
```

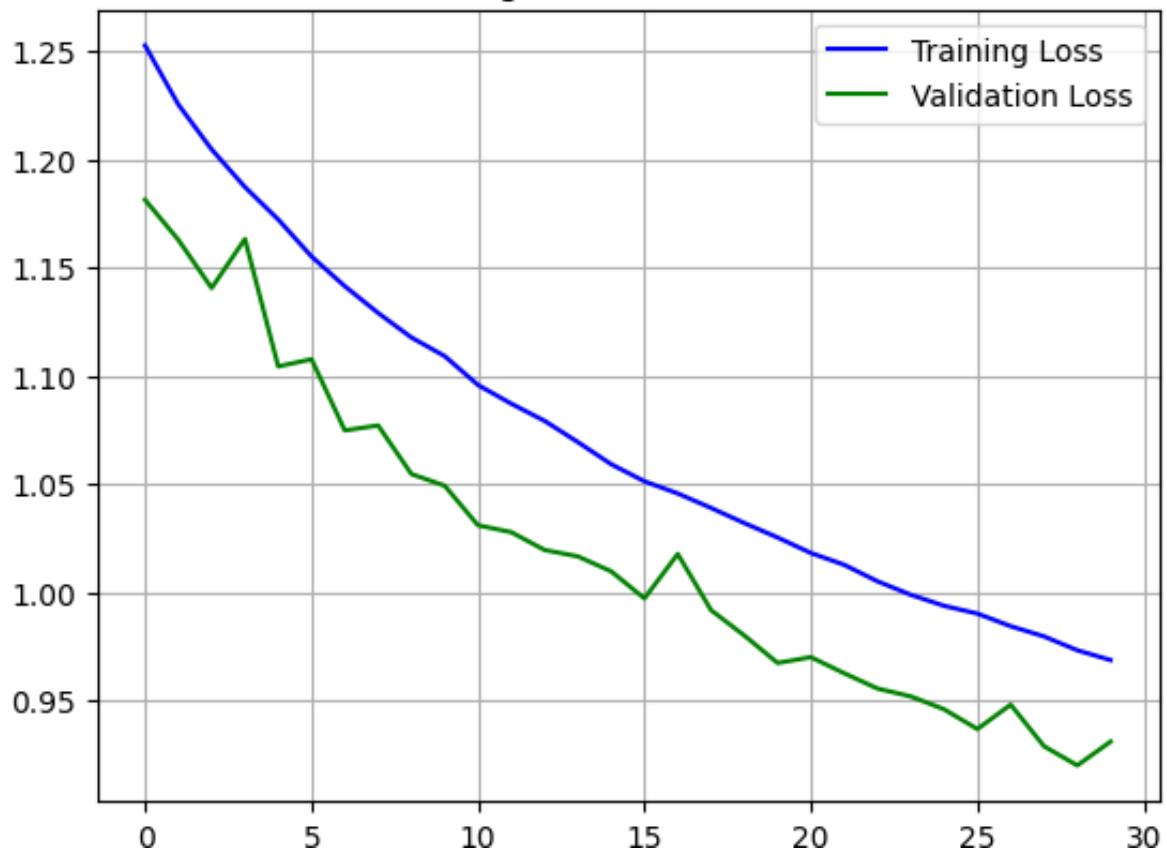
```
In [ ]: epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```



<Figure size 640x480 with 0 Axes>

Training and Validation Loss



```
In [ ]: # Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples)
predicted_classes = np.argmax(predictions, axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {} / {}".format(len(errors), test_generator.samples))
```

<ipython-input-9-1bdf1cd647ce>:14: UserWarning: `Model.predict_generator` is deprecated and will be removed in a future version. Please use `Model.predict`, which supports generators.

predictions = model.predict_generator(test_generator, steps=test_generator.samples/test_generator.batch_size, verbose=1)

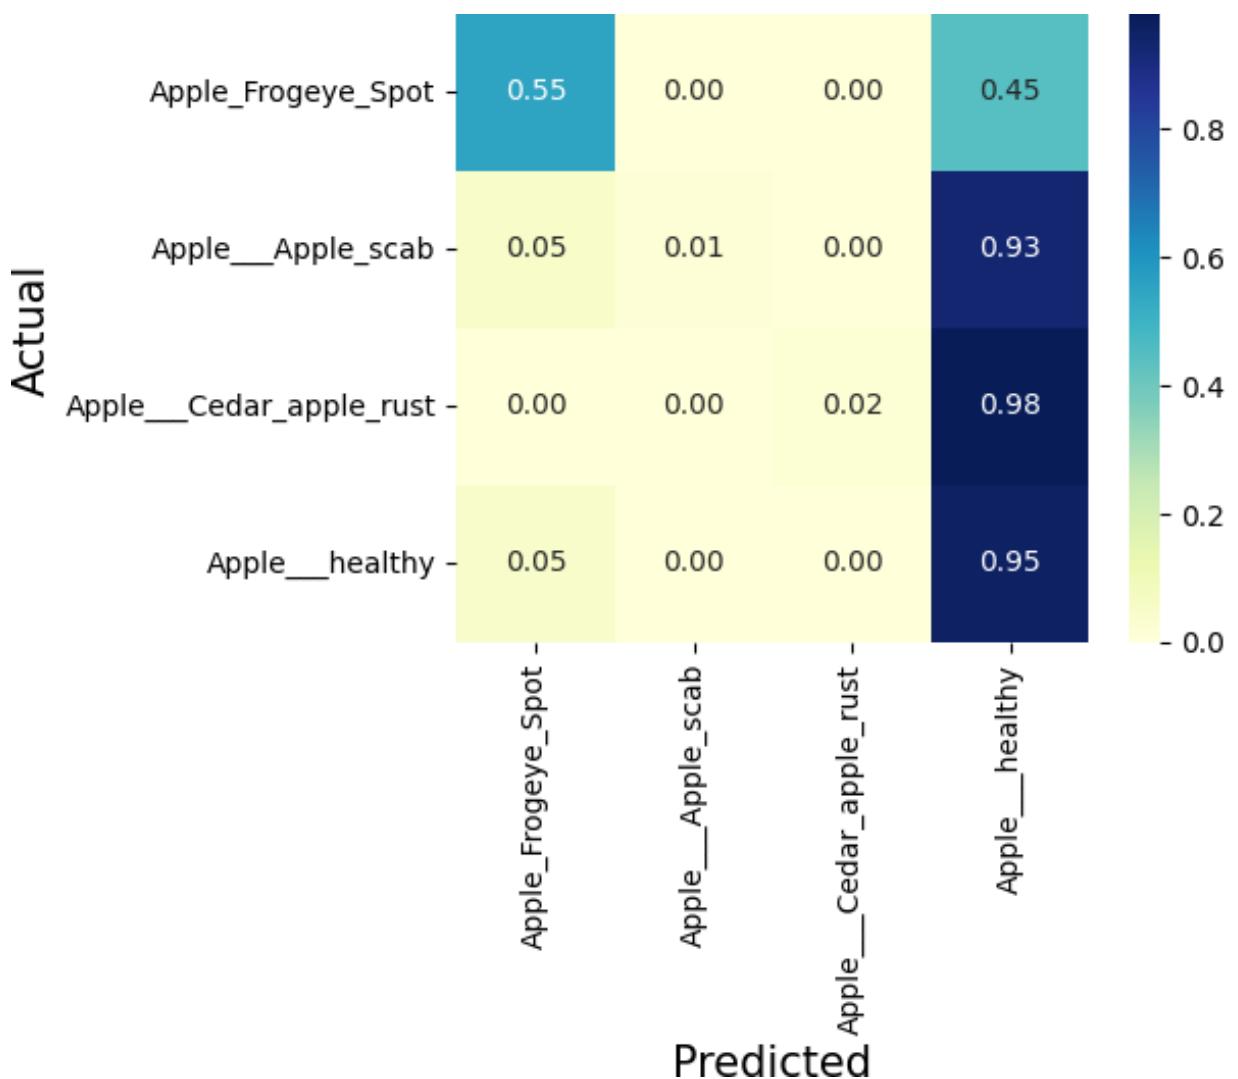
546/546 [=====] - 6s 9ms/step

No of errors = 238/546

```
In [ ]: accuracy = ((test_generator.samples - len(errors)) / test_generator.samples)
accuracy
```

Out[]: 56.41025641025641

```
In [ ]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5, 4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index, yticklabels=label2index,
            plt.ylabel('Actual', fontsize=15),
            plt.xlabel('Predicted', fontsize=15))
plt.show(block=False)
```



```
In [ ]: from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names
```

	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.75	0.55	0.64	103
Apple_Apple_scab	1.00	0.01	0.03	134
Apple_Cedar_apple_rust	1.00	0.02	0.04	49
Apple_healthy	0.53	0.95	0.68	260
accuracy			0.56	546
macro avg	0.82	0.39	0.35	546
weighted avg	0.73	0.56	0.46	546

RESULT: We have successfully implemented an image classification task using pre-trained models like AlexNet, VGGNet, InceptionNet and ResNet and compare the results.

Experiment : 5

AIM: Design a CNN architecture to implement the image classification task over an image dataset. Perform the Hyper-parameter tuning and record the results.

Importing Necessary libraries

- From keras library we are going to use image preprocessing task, to normalize the image pixel values in between 0 to 1.
- Model is imported to load various Neural Network models such as Sequential.
- We are going to use Stochastic Gradient Descent(SGD) as a optimizer
- Keras layers such as Dense, Flatten, Conv2D and MaxPooling is used to implement the CNN model

```
In [ ]: from keras import applications
from keras import optimizers
from keras.models import Sequential, Model
from keras.layers import Dropout, Flatten, Dense, GlobalAveragePooling2D
from keras import backend as k
from keras.callbacks import ModelCheckpoint, LearningRateScheduler, TensorBoard
```

```
In [ ]: import numpy as np
import keras
from keras import models
import matplotlib.pyplot as plt
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Model
from keras.optimizers import SGD
from keras import layers
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
from keras import Input
```

Loading the training and testing data and defining the basic parameters

- We are resizing the input image to 128 * 128
- In the dataset : Training Set : 70% Validation Set : 20% Test Set : 10%

```
In [ ]: # Normalize training and validation data in the range of 0 to 1
train_datagen = ImageDataGenerator(rescale=1./255) # vertical_flip=True,
                                                    # horizontal_flip=True
                                                    # height_shift_range=0
                                                    # width_shift_range=0.
validation_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

# Read the training sample and set the batch size
train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/plant_v1',
    target_size=(64, 64),
    batch_size=16,
    class_mode='categorical')

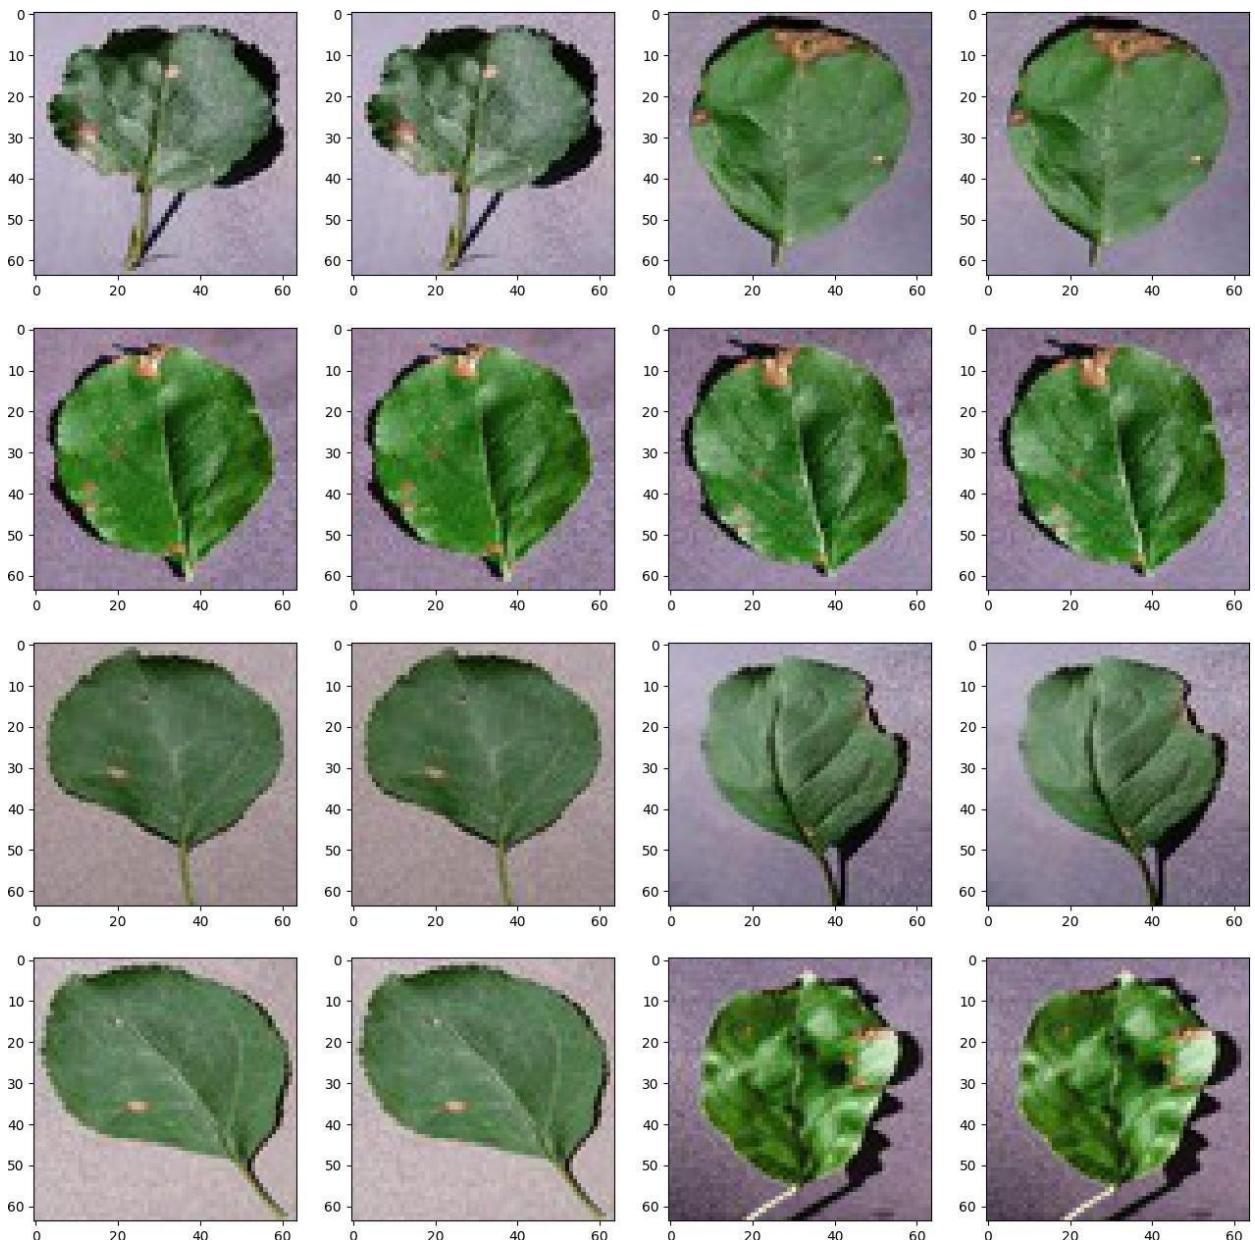
# Read Validation data from directory and define target size with batch s
validation_generator = validation_datagen.flow_from_directory(
    '/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/plant_v1',
    target_size=(64, 64),
    batch_size=16,
    class_mode='categorical',
    shuffle=False)

test_generator = test_datagen.flow_from_directory(
    '/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/plant_v1',
    target_size=(64, 64),
    batch_size=1,
    class_mode='categorical',
    shuffle=False)
```

Found 3002 images belonging to 4 classes.
Found 635 images belonging to 4 classes.
Found 546 images belonging to 4 classes.

Visualization of few images

```
In [ ]: plt.figure(figsize=(16, 16))
for i in range(1, 17):
    plt.subplot(4, 4, i)
    img, label = test_generator.next()
    # print(img.shape)
    # print(label)
    plt.imshow(img[0])
plt.show()
```



```
In [ ]: img, label = test_generator.next()  
img[0].shape
```

```
Out[ ]: (64, 64, 3)
```

Model Definition

- We are going to use 2 convolution layers with 3×3 filter and relu as an activation function
- Then max pooling layer with 2×2 filter is used
- After that we are going to use Flatten layer
- Then Dense layer is used with relu function
- In the output layer softmax function is used with 4 neurons as we have four class dataset.

- `model.summary()` is used to check the overall architecture of the model with number of learnable parameters in each

```
In [ ]: # Create the model
model = models.Sequential()
# Add new layers
model.add(Conv2D(128, kernel_size=(3,3), activation = 'relu', input_shape=(28,28,3)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(layers.Flatten())
model.add(layers.Dense(32, activation='relu'))
model.add(layers.Dense(4, activation='softmax'))
model.summary()
```

Compiling and Training the Model

```
In [ ]: # sgd = SGD(lr=0.001, decay=1e-6, momentum=0.9, nesterov=True)
# We are going to use accuracy metrics and cross entropy loss as performance metrics
model.compile(optimizer = optimizers.Adam(learning_rate = 0.001), loss='categorical_crossentropy')
# Train the model
history = model.fit(train_generator,
                     steps_per_epoch=train_generator.samples/train_generator.batch_size,
                     epochs=30,
                     validation_data=validation_generator,
                     validation_steps=validation_generator.samples/validation_generator.batch_size,
                     verbose=1)
```

Saving the model

```
In [ ]: model.save("CONV_plant_deseas.h5")
print("Saved model to disk")
```

Saved model to disk

Loading the model

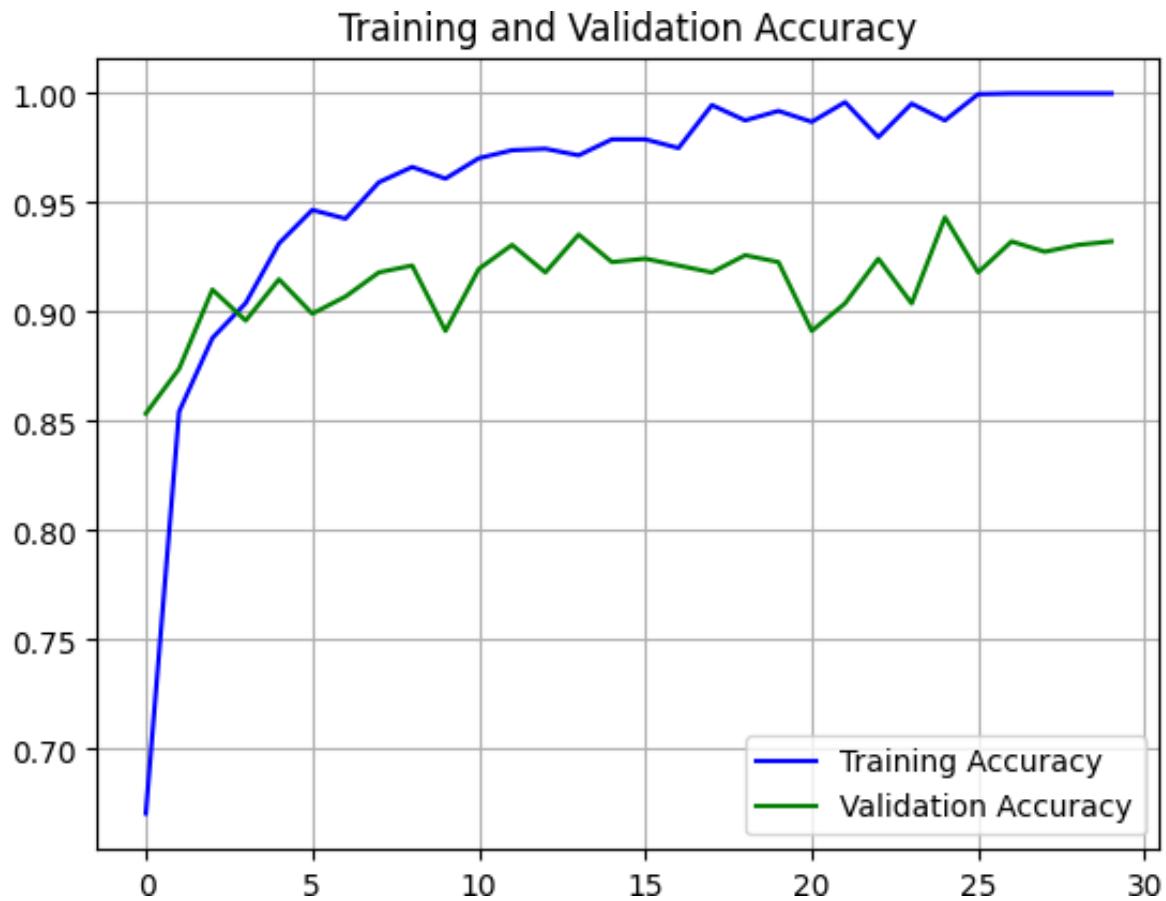
```
In [ ]: model = models.load_model('CONV_plant_deseas.h5')
```

Visualization of Accuracy and Loss Curves

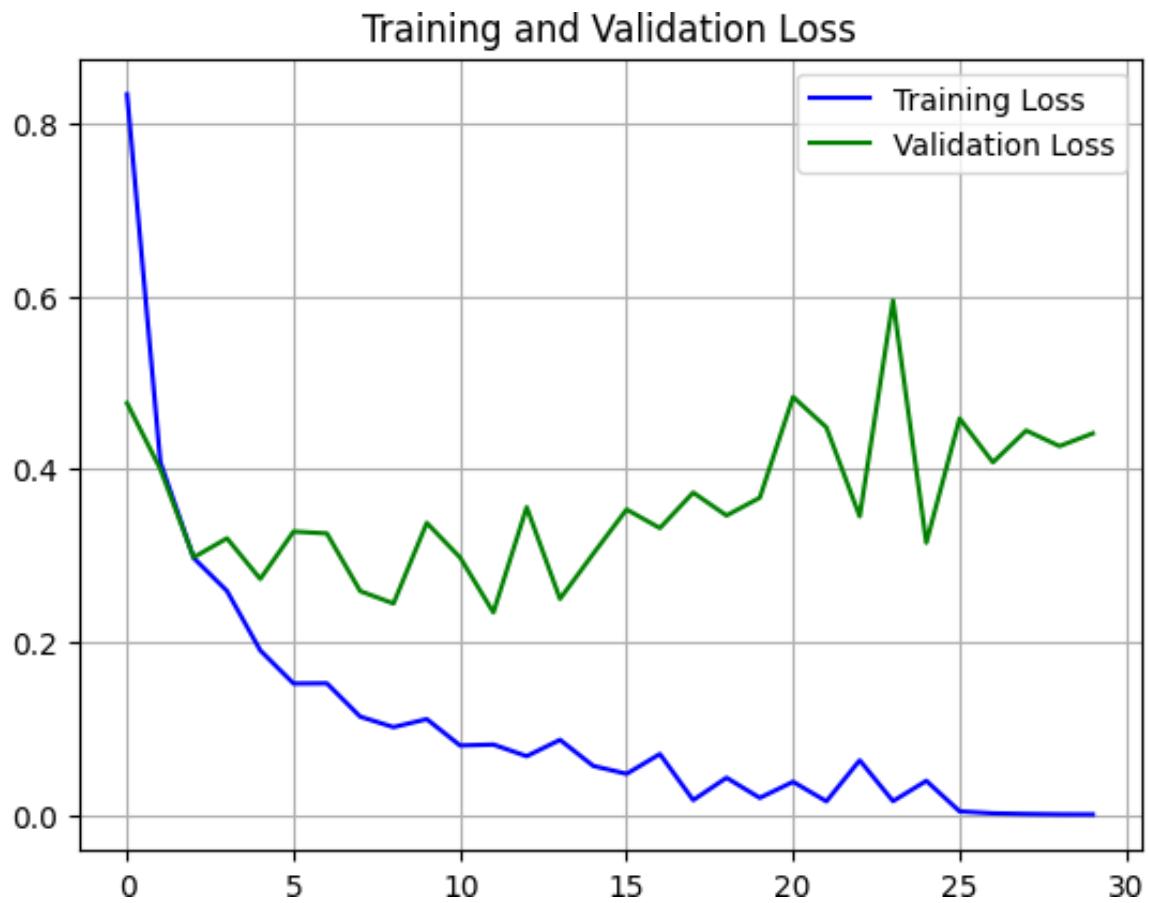
```
In [ ]: train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']
```

```
In [ ]: epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```



<Figure size 640x480 with 0 Axes>



Prediction

```
In [ ]: # Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples)
predicted_classes = np.argmax(predictions, axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}/{}".format(len(errors), test_generator.samples))
```

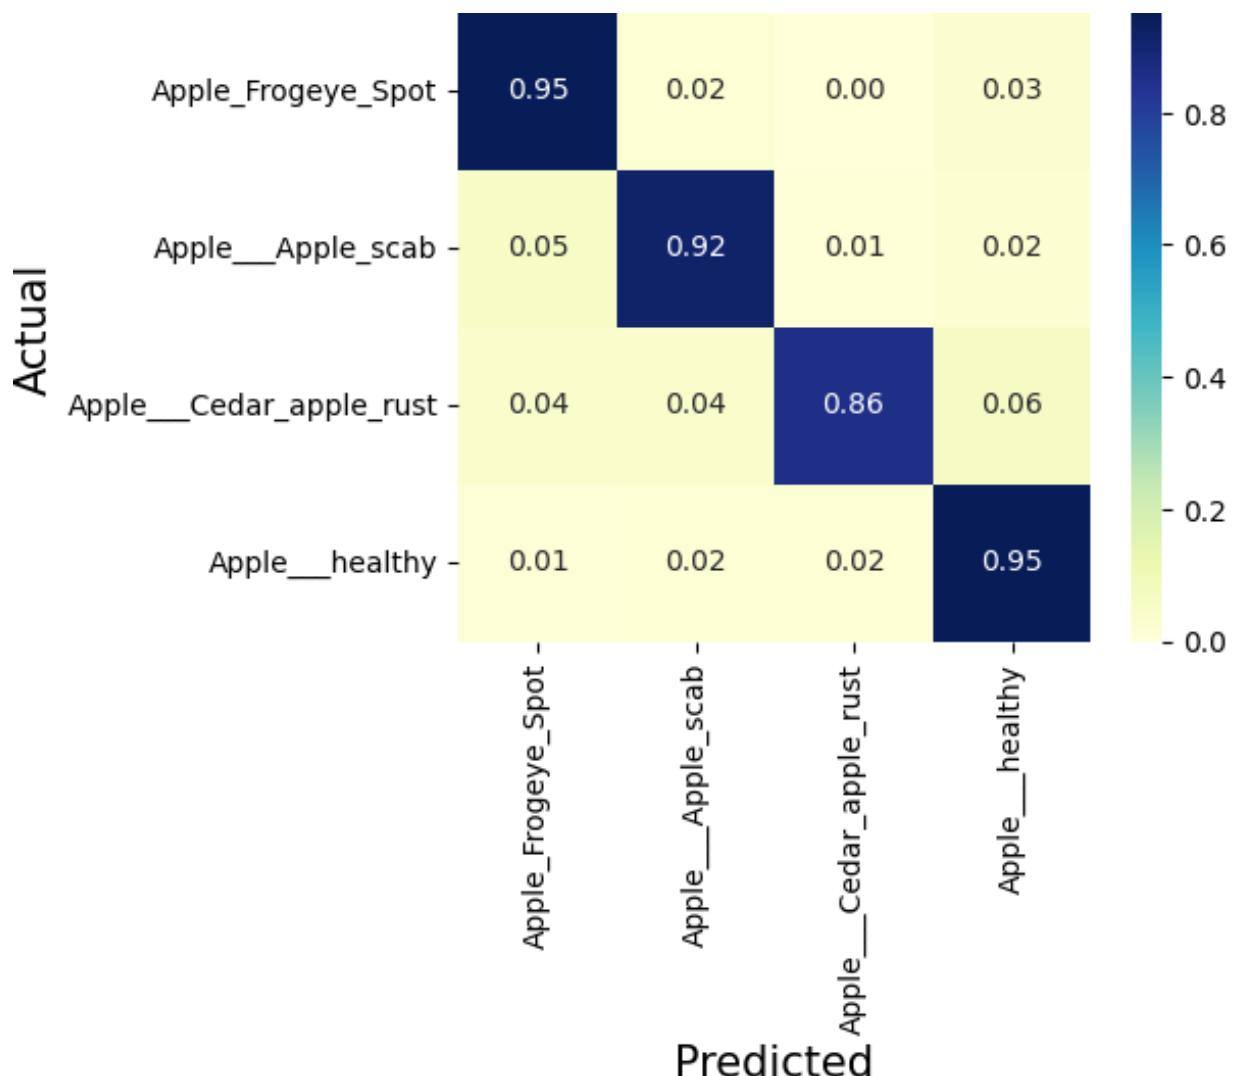
```
<ipython-input-12-1bdf1cd647ce>:14: UserWarning: `Model.predict_generator`  
is deprecated and will be removed in a future version. Please use `Model.p  
redict`, which supports generators.  
    predictions = model.predict_generator(test_generator, steps=test_generat  
or.samples/test_generator.batch_size,verbose=1)  
546/546 [=====] - 249s 456ms/step  
No of errors = 36/546
```

```
In [ ]: accuracy = ((test_generator.samples-len(errors))/test_generator.samples)  
accuracy
```

```
Out[ ]: 93.4065934065934
```

Confusion Matrix

```
In [ ]: from sklearn.metrics import confusion_matrix  
import seaborn as sns  
import numpy as np  
from matplotlib import pyplot as plt  
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)  
cm = np.array(cm)  
# Normalise  
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]  
fig, ax = plt.subplots(figsize=(5,4))  
sns.heatmap(cmn, annot=True, fmt=' .2f', xticklabels=label2index, yticklab  
plt.ylabel('Actual', fontsize=15)  
plt.xlabel('Predicted', fontsize=15)  
plt.show(block=False)
```



Classification Report

```
In [ ]: from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names))

      precision    recall  f1-score   support

Apple_Frogeye_Spot       0.89      0.95      0.92      103
Apple_Apple_scab         0.93      0.92      0.92      134
Apple_Cedar_apple_rust   0.88      0.86      0.87       49
Apple_healthy            0.96      0.95      0.96     260

accuracy                  0.93
macro avg                 0.92      0.92      0.92      546
weighted avg              0.93      0.93      0.93      546
```

```
In [ ]:
```

Finetuning Hyper-parameters

```
In [ ]: def create_model(learn_rate=0.01, momentum=0) :  
    image_size = 128  
    input_shape = (image_size, image_size, 3)  
  
    model = models.Sequential()  
    model.add(Conv2D(128, kernel_size=(3,3), activation = 'relu', input_s  
    model.add(MaxPooling2D(pool_size=(2,2)))  
    model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))  
    model.add(MaxPooling2D(pool_size=(2,2)))  
    model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))  
    model.add(MaxPooling2D(pool_size=(2,2)))  
    model.add(layers.Flatten())  
    model.add(layers.Dense(32, activation='relu'))  
    model.add(layers.Dense(4, activation='softmax'))  
  
    # model = Model(input_shape, x)  
  
    model.compile(loss='categorical_crossentropy',  
                  optimizer=optimizers.SGD(lr=learn_rate, momentum=momen  
                  metrics=['accuracy'])  
  
    return model
```

```
In [ ]: learn_rate = [1e-9, 1e-3]  
momentum = [0.6, 0.9]  
  
def try_fit(learn_rate,momentum) :  
    history_page=[]  
    for lr in learn_rate:  
        for moment in momentum:  
            model = create_model(lr,moment)  
            history = model.fit_generator(  
                train_generator,  
                epochs=1,  
                validation_data=validation_generator)  
            history_page.append(history)  
    return history_page  
  
history_page = try_fit(learn_rate,momentum)  
history_page[0].history['accuracy']
```

```
<ipython-input-26-491e2ac86ce5>:10: UserWarning: `Model.fit_generator` is  
deprecated and will be removed in a future version. Please use `Model.fit  
    history = model.fit_generator(
```

```
188/188 [=====] - 305s 2s/step - loss: 0.6797 - accuracy: 0.4417 - val_loss: 0.6787 - val_accuracy: 0.5039  
188/188 [=====] - 299s 2s/step - loss: 0.7016 - accuracy: 0.2375 - val_loss: 0.7035 - val_accuracy: 0.2000  
188/188 [=====] - 295s 2s/step - loss: 0.5859 - accuracy: 0.3944 - val_loss: 0.4985 - val_accuracy: 0.5181  
188/188 [=====] - 298s 2s/step - loss: 0.5418 - accuracy: 0.4460 - val_loss: 0.4854 - val_accuracy: 0.5181
```

Out[]: [0.44170552492141724]

In []:

Experiment : 7

AIM: Train a Recurrent Neural Network (RNN) on the IMDB large movie review dataset for sentiment analysis.

Importing Necessary libraries

```
In [ ]: import pandas as pd      # to load dataset
import numpy as np       # for mathematic equation
from nltk.corpus import stopwords    # to get collection of stopwords
from sklearn.model_selection import train_test_split      # for splitting
from tensorflow.keras.preprocessing.text import Tokenizer  # to encode te
from tensorflow.keras.preprocessing.sequence import pad_sequences   # to
from tensorflow.keras.models import Sequential      # the model
from tensorflow.keras.layers import Embedding, LSTM, Dense # layers of th
from tensorflow.keras.callbacks import ModelCheckpoint    # save model
from tensorflow.keras.models import load_model     # load saved model
import re
from keras.layers import SimpleRNN
```

Preparing the data named IMDB

```
In [ ]: data = pd.read_csv('/content/drive/MyDrive/AMITY/Deep Learning (codes)/Da
print(data)
```

Stop Word is a commonly used words in a sentence, usually a search engine is programmed to ignore this words (i.e. "the", "a", "an", "of", etc.) Declaring the english stop words

```
In [ ]: import nltk
nltk.download("stopwords")
english_stops = set(stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]  Unzipping corpora/stopwords.zip.
```

Load and Clean Dataset

In the original dataset, the reviews are still dirty. There are still html tags,

numbers, uppercase, and punctuations. This will not be good for training, so in `load_dataset()` function, beside loading the dataset using pandas, I also pre-process the reviews by removing html tags, non alphabet (punctuations and numbers), stop words, and lower case all of the reviews.

Encode Sentiments

In the same function, We also encode the sentiments into integers (0 and 1). Where 0 is for negative sentiments and 1 is for positive sentiments.

```
In [ ]: def load_dataset():
    df = pd.read_csv('/content/drive/MyDrive/AMITY/Deep Learning (codes)/
        x_data = df['review']           # Reviews/Input
        y_data = df['sentiment']        # Sentiment/Output

        # PRE-PROCESS REVIEW
        x_data = x_data.replace({'<.*?>': ''}, regex = True)          # remove
        x_data = x_data.replace({'[^A-Za-z]': ' '}, regex = True)          # remove
        x_data = x_data.apply(lambda review: [w for w in review.split() if w])
        x_data = x_data.apply(lambda review: [w.lower() for w in review])   # lower

        # ENCODE SENTIMENT -> 0 & 1
        y_data = y_data.replace('positive', 1)
        y_data = y_data.replace('negative', 0)

    return x_data, y_data

x_data, y_data = load_dataset()

print('Reviews')
print(x_data, '\n')
print('Sentiment')
print(y_data)
```

Split Dataset

In this work, We decided to split the data into 80% of Training and 20% of Testing set using `train_test_split` method from Scikit-Learn. By using this method, it automatically shuffles the dataset. We need to shuffle the data because in the original dataset, the reviews and sentiments are in order, where they list positive reviews first and then negative reviews. By shuffling the data, it will be distributed equally in the model, so it will be more accurate for predictions.

```
In [ ]: x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_
```

```

print('Train Set')
print(x_train, '\n')
print(x_test, '\n')
print('Test Set')
print(y_train, '\n')
print(y_test)

```

Function for getting the average review length, by calculating the mean of all the reviews length (using numpy.mean)

```

In [ ]: def get_max_length():
    review_length = []
    for review in x_train:
        review_length.append(len(review))

    return int(np.ceil(np.mean(review_length)))

```

Tokenize and Pad/Truncate Reviews

A Neural Network only accepts numeric data, so we need to encode the reviews. I use tensorflow.keras.preprocessing.text.Tokenizer to encode the reviews into integers, where each unique word is automatically indexed (using fit_on_texts method) based on x_train.

x_train and x_test is converted into integers using texts_to_sequences method.

Each reviews has a different length, so we need to add padding (by adding 0) or truncating the words to the same length (in this case, it is the mean of all reviews length) using

tensorflow.keras.preprocessing.sequence.pad_sequences.

```

In [ ]: # ENCODE REVIEW
token = Tokenizer(lower=False)      # no need lower, because already lowere
token.fit_on_texts(x_train)
x_train = token.texts_to_sequences(x_train)
x_test = token.texts_to_sequences(x_test)

max_length = get_max_length()

x_train = pad_sequences(x_train, maxlen=max_length, padding='post', trunc
x_test = pad_sequences(x_test, maxlen=max_length, padding='post', truncat

total_words = len(token.word_index) + 1    # add 1 because of 0 padding
print('Total Words:', total_words)

print('Encoded X Train\n', x_train, '\n')
print('Encoded X Test\n', x_test, '\n')

```

```
print('Maximum review length: ', max_length)
```

Build Architecture/Model

Embedding Layer: in simple terms, it creates word vectors of each word in the word_index and group words that are related or have similar meaning by analyzing other words around them.

RNN Layer: to make a decision to keep or throw away data by considering the current input, previous output.

Dense Layer: compute the input with the weight matrix and bias (optional), and using an activation function. I use Sigmoid activation function for this work because the output is only 0 or 1.

The optimizer is Adam and the loss function is Binary Crossentropy because again the output is only 0 and 1, which is a binary number.

```
In [ ]: rnn = Sequential()

rnn.add(Embedding(total_words, 32, input_length=max_length))
rnn.add(SimpleRNN(64, input_shape=(total_words, max_length), return_sequences=True))
rnn.add(Dense(1, activation='sigmoid')) #flatten

print(rnn.summary())
rnn.compile(loss="binary_crossentropy", optimizer='adam', metrics=["accuracy"])
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
embedding (Embedding)	(None, 130, 32)	2964352
simple_rnn (SimpleRNN)	(None, 64)	6208
dense (Dense)	(None, 1)	65
<hr/>		
Total params: 2,970,625		
Trainable params: 2,970,625		
Non-trainable params: 0		

None

Trainin the Model

```
In [ ]: history = rnn.fit(x_train, y_train, epochs=20, batch_size=128, verbose=1)
```

Saving The Model

```
In [ ]: model = rnn.save('rnn.h5')
loaded_model = load_model('rnn.h5')
```

Evaluation

```
In [ ]: y_pred = rnn.predict(x_test, batch_size = 128)
print(y_pred)
print(y_test)
for i in range(len(y_pred)):
    if y_pred[i]>0.5:
        y_pred[i] = 1
    else:
        y_pred[i] = 0

true = 0
for i, y in enumerate(y_test):
    if y == y_pred[i]:
        true += 1

print('Correct Prediction: {}'.format(true))
print('Wrong Prediction: {}'.format(len(y_pred) - true))
print('Accuracy: {}'.format(true/len(y_pred)*100))
```

Message: Nothing was typical about this. Everything was beautifully done in this movie, the story, the flow, the scenario, everything. I highly recommend it for mystery lovers, for anyone who wants to watch a good movie!

Example review

```
In [ ]: review = str(input('Movie Review: '))
```

Movie Review: Nothing was typical about this. Everything was beautifully done in this movie, the story, the flow, the scenario, everything. I highly recommend it for mystery lovers, for anyone who wants to watch a good movie!

Pre-processing of entered review

```
In [ ]: # Pre-process input
regex = re.compile(r'[^a-zA-Z\s]')
review = regex.sub(' ', review)
print('Cleaned: ', review)
```

```

words = review.split(' ')
filtered = [w for w in words if w not in english_stops]
filtered = ' '.join(filtered)
filtered = [filtered.lower()]

print('Filtered: ', filtered)

```

Cleaned: Nothing was typical about this Everything was beautifully done in this movie the story the flow the scenario everything I highly recommend it for mystery lovers for anyone who wants to watch a good movie
 Filtered: ['nothing typical everything beautifully done movie story flow scenario everything i highly recommend mystery lovers anyone wants watch g ood movie']

```
In [ ]: tokenize_words = token.texts_to_sequences(filtered)
tokenize_words = pad_sequences(tokenize_words, maxlen=max_length, padding='post')
print(tokenize_words)
```

```
[[ 76  705  174 1210  126   3   13 2692 2596  174    1  442  280  701
  1771 155  400   33    9   3   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0]]
```

Prediction

```
In [ ]: result = rnn.predict(tokenize_words)
print(result)
```

```
1/1 [=====] - 0s 40ms/step
[[0.78446704]]
```

```
In [ ]: if result >= 0.7:
        print('positive')
    else:
        print('negative')
```

```
positive
```

RESULT: We have successfully trained a Recurrent Neural Network (RNN) on the IMDB large movie review dataset for sentiment analysis.

Experiment : 8

AIM: Download the tweets data from twitter and run an RNN/LSTM to classify the sentiment as Hate/ Non-hate speech. Compare the results of both the models.

Importing Necessary libraries

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Preparing the data named IMDB

```
In [ ]: df = pd.read_csv('/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data'
df.head()
```

	Unnamed:	count	hate_speech	offensive_language	neither	class			
0	0	3	0	0	3	2	@mayasol As a woma shoul		
1	1	3	0	3	0	1	@mleew1 dats cold. dw	!	
2	2	3	0	3	0	1	@UrKindOf Dawg @80sb	!!!	
3	3	3	0	2	1	1	@C_G_And @viva_base	!!!!	
4	4	6	0	6	0	1	@ShenikaRo The shit	!!!!!!	

```
In [ ]: classes = ['Hate Speech', 'Offensive Language', 'None']
```

```
In [ ]: df.drop(['count', 'hate_speech', 'offensive_language', 'neither', 'Unnamed: 0']
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	class	tweet
0	2	!!! RT @mayasolovely: As a woman you shouldn't...
1	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	1	!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	1	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...

```
In [ ]: df.shape
```

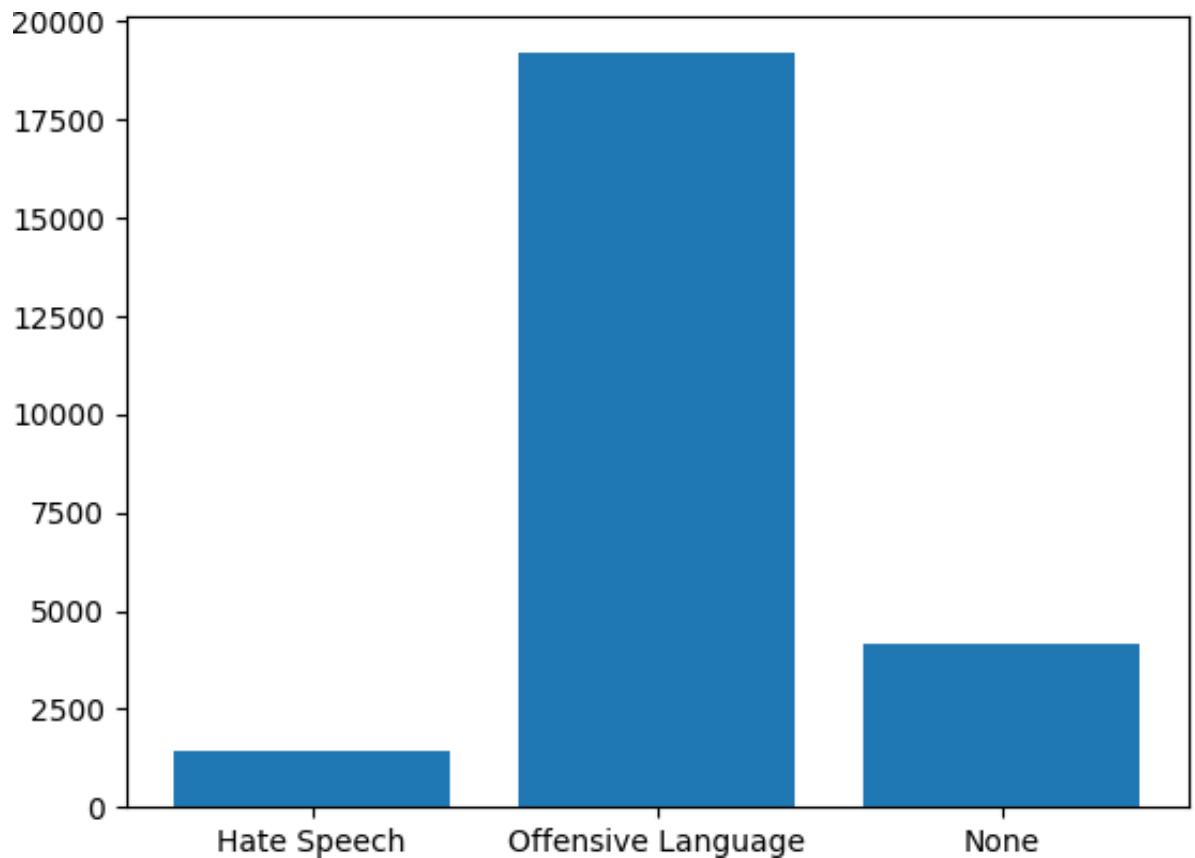
```
Out[ ]: (24783, 2)
```

Statistics of the Data

```
In [ ]:
```

```
labels = df['class']
unique, counts = np.unique(labels, return_counts=True)
values = list(zip(unique, counts))
plt.bar(classes, counts)
for i in values:
    print(classes[i[0]], ' : ', i[1])
plt.show()
```

```
Hate Speech : 1430
Offensive Language : 19190
None : 4163
```



```
In [ ]: hate_tweets = df[df['class']==0]
        offensive_tweets = df[df['class']==1]
        neither = df[df['class']==2]
        print(hate_tweets.shape)
        print(offensive_tweets.shape)
        print(neither.shape)
```

```
(1430, 2)
(19190, 2)
(4163, 2)
```

```
In [ ]: for i in range(3):
        hate_tweets = pd.concat([hate_tweets,hate_tweets],ignore_index = True
        neither = pd.concat([neither,neither,neither], ignore_index = True)
        offensive_tweets = offensive_tweets.iloc[0:12000,:]
        print(hate_tweets.shape)
        print(offensive_tweets.shape)
        print(neither.shape)
```

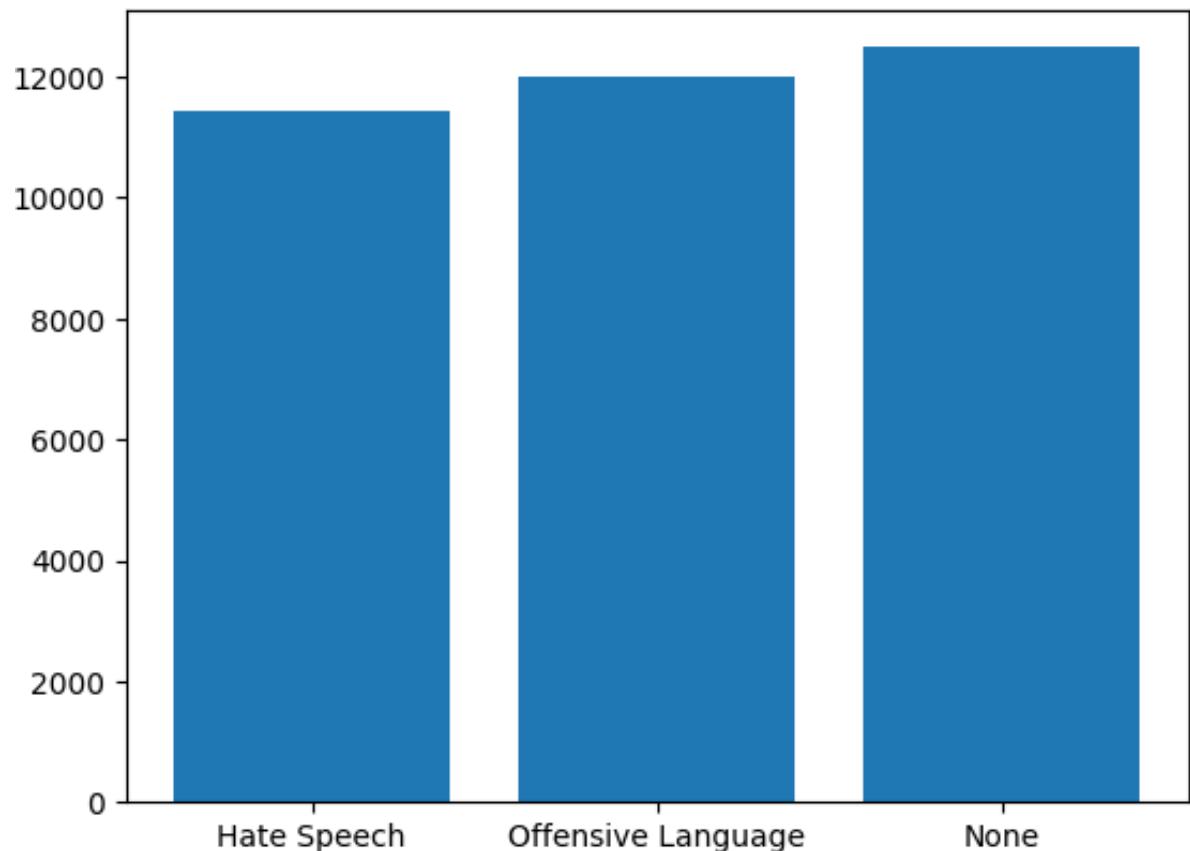
```
(11440, 2)
(12000, 2)
(12489, 2)
```

```
In [ ]: df = pd.concat([hate_tweets,offensive_tweets,neither],ignore_index = True
df.shape
```

```
Out[ ]: (35929, 2)
```

```
In [ ]: labels = df['class']
unique, counts = np.unique(labels, return_counts=True)
values = list(zip(unique, counts))
plt.bar(classes,counts)
for i in values:
    print(classes[i[0]],' : ',i[1])
plt.show()
```

Hate Speech : 11440
Offensive Language : 12000
None : 12489



```
In [ ]: df.head()
```

	class	tweet
0	0	"@Blackman38Tide: @WhaleLookyHere @HowdyDowdy1...
1	0	"@CB_Baby24: @white_thunduh alsarabsss" hes a ...
2	0	"@DevilGrimz: @VigxRArts you're fucking gay, b...
3	0	"@MarkRoundtreeJr: LMFAOOOO I HATE BLACK PEOPL...
4	0	"@NoChillPaz: "At least I'm not a nigger" http...

Preprocessing

```
In [ ]: import nltk
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
import re
nltk.download('wordnet')
nltk.download('stopwords')
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]  Unzipping corpora/stopwords.zip.
```

Out[]: True

```
In [ ]: # dealing with Slangs
d = {'luv':'love', 'wud':'would', 'lyk':'like', 'wateva':'whatever', 'ttyp':'kul':'cool', 'fyn':'fine', 'omg':'oh my god!', 'fam':'family', 'cud':'could', 'fud':'food', 'u':'you', 'ur':'your', 'bday' : 'birthday', 'bihday' : 'birthday'}
```

```
In [ ]: stop_words = set(stopwords.words("english"))
stop_words.add('rt')
stop_words.remove('not')
lemmatizer = WordNetLemmatizer()
giant_url_regex = ('http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@.&+])|' '![!*\(\)\],|('
mention_regex = '@[\w\-\]+'

def clean_text(text):
    text = re.sub('\"', "", text)
    text = re.sub(mention_regex, ' ', text) #removing all user names
    text = re.sub(giant_url_regex, ' ', text) #removing the urls
    text = text.lower()
    text = re.sub("hmm+", "", text) #removing variants of hmmm
    text = re.sub("[^a-z]+", " ", text) #removing all numbers, special ch
    text = text.split()
    text = [word for word in text if not word in stop_words]
    text = [d[word] if word in d else word for word in text] #replacing
    text = [lemmatizer.lemmatize(token) for token in text]
    text = [lemmatizer.lemmatize(token, "v") for token in text]
    text = " ".join(text)
    return text
```

```
In [ ]: df['processed_tweets'] = df.tweet.apply(lambda x: clean_text(x)) # df.r
df.head()
```

Out[]:	class	tweet	processed_tweets
0	0	"@Blackman38Tide: @WhaleLookyHere @HowdyDowdy1...	queer gaywad
1	0	"@CB_Baby24: @white_thunduh alsarabsss" hes a ...	alsarabsss he beaner smh tell he mexican
2	0	"@DevilGrimz: @VigxRArts you're fucking gay, b...	fuck gay blacklist hoe hold tehgodclan anyway
3	0	"@MarkRoundtreeJr: LMFAOOOO I HATE BLACK PEOPL...	lmfaoooo hate black people black people nigger
4	0	"@NoChillPaz: "At least I'm not a nigger" http...	least not nigger lmfao

```
In [ ]: x = df['processed_tweets']
y = df['class']
print(x.shape)
print(y.shape)
```

(35929,)
(35929,)

```
In [ ]: # finding unique words
word_unique = []
for i in x:
    for j in i.split():
        word_unique.append(j)
unique, counts = np.unique(word_unique, return_counts=True)
print("The total words in the tweets are : ", len(word_unique))
print("The total UNIQUE words in the tweets are : ", len(unique))
```

The total words in the tweets are : 275540
The total UNIQUE words in the tweets are : 14146

```
In [ ]: # finding length of tweets
tweets_length = []
for i in x:
    tweets_length.append(len(i.split()))
print("The Average Length tweets are : ", np.mean(tweets_length))
print("The max length of tweets is : ", np.max(tweets_length))
print("The min length of tweets is : ", np.min(tweets_length))
```

The Average Length tweets are : 7.669013888502324
The max length of tweets is : 28
The min length of tweets is : 0

```
In [ ]: tweets_length = pd.DataFrame(tweets_length)
# tweets_length.describe()
```

Out[]:

count	35929.000000
mean	7.669014
std	3.989625
min	0.000000
25%	4.000000
50%	7.000000
75%	11.000000
max	28.000000

In []:

```
#Sorting the Unique words based on their Frequency
col = list(zip(unique, counts))
col = sorted(col, key = lambda x: x[1], reverse=True)
col=pd.DataFrame(col)
print("Top 20 Occuring Words with their frequency are:")
col.iloc[:20,:]
```

Top 20 Occuring Words with their frequency are:

Out[]:

		0	1
0	bitch	9066	
1	like	3817	
2	get	3636	
3	hoe	3426	
4	trash	3217	
5	fuck	3103	
6	nigga	2819	
7	faggot	2239	
8	as	2073	
9	you	1851	
10	pussy	1847	
11	go	1773	
12	not	1764	
13	bird	1515	
14	lol	1494	
15	nigger	1459	
16	say	1456	
17	make	1373	
18	amp	1329	
19	white	1328	

```
In [ ]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [ ]: vectorizer = TfidfVectorizer(max_features = 8000 )
# tokenize and build vocab

vectorizer.fit(x)
# summarize

print(len(vectorizer.vocabulary_))
print(vectorizer.idf_.shape)
```

8000
(8000,)

```
In [ ]: x_tfidf = vectorizer.transform(x).toarray()
print(x_tfidf.shape)
```

```
(35929, 8000)
```

```
In [ ]: from keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
In [ ]: num_words = 8000
embed_dim = 32
tokenizer = Tokenizer(num_words=num_words, oov_token = "<oov>")
tokenizer.fit_on_texts(x)
word_index=tokenizer.word_index
sequences = tokenizer.texts_to_sequences(x)
length=[]
for i in sequences:
    length.append(len(i))
print(len(length))
print("Mean is: ", np.mean(length))
print("Max is: ", np.max(length))
print("Min is: ", np.min(length))
```

```
35929
```

```
Mean is:  7.669013888502324
Max is:  28
Min is:  0
```

```
In [ ]: pad_length = 24
sequences = pad_sequences(sequences, maxlen = pad_length, truncating = 'pre')
sequences.shape
```

```
Out[ ]: (35929, 24)
```

Splitting the Data

```
In [ ]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(sequences,y,test_size = 0.2)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(34132, 24)
(1797, 24)
(34132,)
(1797,)
```

RNN Model

```
In [ ]: from keras.layers import Dense, Embedding, Dropout, Activation, Flatten,
from keras.layers import GlobalMaxPool1D
from keras.models import Model, Sequential
```

```
import tensorflow as tf
```

```
In [ ]: recall = tf.keras.metrics.Recall()
precision = tf.keras.metrics.Precision()

model = Sequential([
    Embedding(num_words, embed_dim, input_length = pad_length),
    SimpleRNN(8, return_sequences = True),
    GlobalMaxPooling1D(),
    Dense(20, activation = 'relu', kernel_initializer='he_uniform'),
    Dropout(0.25),
    Dense(3, activation = 'softmax')])
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam',
# model.name = 'Twitter Hate Text Classification'
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
embedding_1 (Embedding)	(None, 24, 32)	256000
simple_rnn_1 (SimpleRNN)	(None, 24, 8)	328
global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 8)	0
dense_2 (Dense)	(None, 20)	180
dropout_1 (Dropout)	(None, 20)	0
dense_3 (Dense)	(None, 3)	63
<hr/>		
Total params: 256,571		
Trainable params: 256,571		
Non-trainable params: 0		

Training the Model

```
In [ ]: history = model.fit(x = x_train, y = y_train, epochs = 5, validation_split
```

```

Epoch 1/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.2029
- accuracy: 0.9373 - val_loss: 0.1583 - val_accuracy: 0.9473
Epoch 2/5
1014/1014 [=====] - 11s 11ms/step - loss: 0.1271
- accuracy: 0.9624 - val_loss: 0.1521 - val_accuracy: 0.9467
Epoch 3/5
1014/1014 [=====] - 12s 11ms/step - loss: 0.0906
- accuracy: 0.9754 - val_loss: 0.1479 - val_accuracy: 0.9537
Epoch 4/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.0714
- accuracy: 0.9801 - val_loss: 0.1644 - val_accuracy: 0.9525
Epoch 5/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.0588
- accuracy: 0.9839 - val_loss: 0.1541 - val_accuracy: 0.9584

```

Evaluation

```
In [ ]: evaluate = model.evaluate(x_test,y_test)
57/57 [=====] - 0s 3ms/step - loss: 0.1112 - accuracy: 0.9688
```

```
In [ ]: print("Test Accuracy is : {:.2f} %".format(evaluate[1]*100))
print("Test Loss is : {:.4f}".format(evaluate[0]))
```

```
Test Accuracy is : 96.88 %
Test Loss is : 0.1112
```

```
In [ ]: predictions = model.predict(x_test)
57/57 [=====] - 0s 2ms/step
```

```
In [ ]: predict = []
for i in predictions:
    predict.append(np.argmax(i))
```

Confusion Matrix

```
In [ ]: from sklearn import metrics
cm = metrics.confusion_matrix(predict,y_test)
acc = metrics.accuracy_score(predict,y_test)
```

```
In [ ]: print("The Confusion matrix is: \n",cm)
```

```
The Confusion matrix is:
[[548  22   1]
 [  6 572   8]
 [  0  19 621]]
```

Accuracy

```
In [ ]: print(acc*100)
```

96.88369504730106

Classification Report

```
In [ ]: from sklearn import metrics
print(metrics.classification_report(y_test, predict))
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	554
1	0.98	0.93	0.95	613
2	0.97	0.99	0.98	630
accuracy			0.97	1797
macro avg	0.97	0.97	0.97	1797
weighted avg	0.97	0.97	0.97	1797

LSTM

```
In [ ]: from tensorflow.keras.layers import Embedding, LSTM, Dense
# ARCHITECTURE
EMBED_DIM = 32
LSTM_OUT = 64

model = Sequential()
model.add(Embedding(num_words, EMBED_DIM, input_length = pad_length))
model.add(LSTM(LSTM_OUT))
model.add(Dense(3, activation='softmax'))
model.compile(optimizer = 'adam', loss = 'sparse_categorical_crossentropy')

print(model.summary())
```

```
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
<hr/>		
embedding_4 (Embedding)	(None, 24, 32)	256000
lstm_1 (LSTM)	(None, 64)	24832
dense_5 (Dense)	(None, 3)	195
<hr/>		
Total params: 281,027		
Trainable params: 281,027		
Non-trainable params: 0		

```
None
```

Training Model

```
In [ ]: history = model.fit(x = x_train, y = y_train, epochs = 10, validation_spli
```

```
Epoch 1/10
1014/1014 [=====] - 26s 24ms/step - loss: 0.4328
- accuracy: 0.8246 - val_loss: 0.2531 - val_accuracy: 0.9127
Epoch 2/10
1014/1014 [=====] - 22s 22ms/step - loss: 0.1874
- accuracy: 0.9405 - val_loss: 0.1876 - val_accuracy: 0.9367
Epoch 3/10
1014/1014 [=====] - 23s 23ms/step - loss: 0.1426
- accuracy: 0.9554 - val_loss: 0.1991 - val_accuracy: 0.9391
Epoch 4/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.1188
- accuracy: 0.9626 - val_loss: 0.1991 - val_accuracy: 0.9361
Epoch 5/10
1014/1014 [=====] - 23s 23ms/step - loss: 0.0999
- accuracy: 0.9688 - val_loss: 0.2142 - val_accuracy: 0.9408
Epoch 6/10
1014/1014 [=====] - 23s 22ms/step - loss: 0.0835
- accuracy: 0.9730 - val_loss: 0.1820 - val_accuracy: 0.9508
Epoch 7/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.0735
- accuracy: 0.9767 - val_loss: 0.2202 - val_accuracy: 0.9473
Epoch 8/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.0670
- accuracy: 0.9783 - val_loss: 0.2666 - val_accuracy: 0.9461
Epoch 9/10
1014/1014 [=====] - 22s 22ms/step - loss: 0.0577
- accuracy: 0.9815 - val_loss: 0.1988 - val_accuracy: 0.9490
Epoch 10/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.0537
- accuracy: 0.9823 - val_loss: 0.2313 - val_accuracy: 0.9490
```

Evaluation

```
In [ ]: evaluate = model.evaluate(x_test,y_test)
57/57 [=====] - 0s 8ms/step - loss: 0.1719 - accuracy: 0.9594

In [ ]: print("Test Accuracy is : {:.2f} %".format(evaluate[1]*100))
print("Test Loss is : {:.4f}".format(evaluate[0]))

Test Accuracy is : 95.94 %
Test Loss is : 0.1719

In [ ]: predictions = model.predict(x_test)

57/57 [=====] - 1s 8ms/step

In [ ]: predict = []
for i in predictions:
    predict.append(np.argmax(i))
```

Confusion Matrix

```
In [ ]: from sklearn import metrics
cm = metrics.confusion_matrix(predict,y_test)
acc = metrics.accuracy_score(predict,y_test)

In [ ]: print("The Confusion matrix is: \n",cm)

The Confusion matrix is:
[[552  42   1]
 [  0 543   0]
 [  2  28 629]]
```

Accuracy

```
In [ ]: print(acc*100)

95.93767390094602
```

Classification Report

```
In [ ]: from sklearn import metrics
print(metrics.classification_report(y_test, predict))
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	554
1	1.00	0.89	0.94	613
2	0.95	1.00	0.98	630
accuracy			0.96	1797
macro avg	0.96	0.96	0.96	1797
weighted avg	0.96	0.96	0.96	1797

In []:

Experiment:-9

AIM:- Implement GAN architecture on MNIST dataset to recognize the handwritten digits.

✓ Importing Necessary libraries

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model
```

✓ Pre-processing

Normalizing the data into 0 to 1 and reshaping the size

```
def preprocess(array):
    """
    Normalizes the supplied array and reshapes it into the appropriate format.
    """

    array = array.astype("float32") / 255.0
    array = np.reshape(array, (len(array), 28, 28, 1))
    return array
```

✓ Adding noise to the original images

```
def noise(array):
    """
    Adds random noise to each image in the supplied array.

    noise_factor = 0.4
    noisy_array = array + noise_factor * np.random.normal(
        loc=0.0, scale=1.0, size=array.shape
    )

    return np.clip(noisy_array, 0.0, 1.0)
```

✓ Visualizing the images

```
def display(array1, array2):
    """
    Displays ten random images from each one of the supplied arrays.
    """

    n = 10

    indices = np.random.randint(len(array1), size=n)
    images1 = array1[indices, :]
    images2 = array2[indices, :]

    plt.figure(figsize=(20, 4))
    for i, (image1, image2) in enumerate(zip(images1, images2)):
        ax = plt.subplot(2, n, i + 1)
        plt.imshow(image1.reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

        ax = plt.subplot(2, n, i + 1 + n)
        plt.imshow(image2.reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

https://colab.research.google.com/drive/1d9ZpfxsVIBC5sjvfR9jSkYXPGI0UePwN#printMode=true
```

plt.show()

✓ Preparing the data

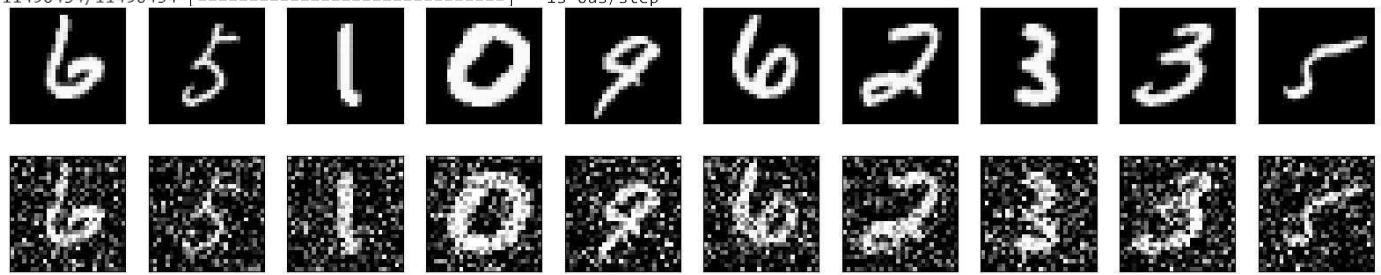
```
# Since we only need images from the dataset to encode and decode, we
# won't use the labels.
(train_data, _), (test_data, _) = mnist.load_data()

# Normalize and reshape the data
train_data = preprocess(train_data)
test_data = preprocess(test_data)

# Create a copy of the data with added noise
noisy_train_data = noise(train_data)
noisy_test_data = noise(test_data)

# Display the train data and a version of it with added noise
display(train_data, noisy_train_data)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 1s 0us/step



✓ Building the Autoencoder

```
input = layers.Input(shape=(28, 28, 1))

# Encoder
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(input)
x = layers.MaxPooling2D((2, 2), padding="same")(x)
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(x)
x = layers.MaxPooling2D((2, 2), padding="same")(x)

# Decoder
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu", padding="same")(x)
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu", padding="same")(x)
x = layers.Conv2D(1, (3, 3), activation="sigmoid", padding="same")(x)

# Autoencoder
autoencoder = Model(input, x)
autoencoder.compile(optimizer="adam", loss="binary_crossentropy")
autoencoder.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[None, 28, 28, 1]	0
conv2d (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	9248

```

max_pooling2d_1 (MaxPooling2D)          0
conv2d_transpose (Conv2DTranspose)       9248
conv2d_transpose_1 (Conv2DTranspose)     9248
conv2d_2 (Conv2D)                      289
=====
Total params: 28,353
Trainable params: 28,353
Non-trainable params: 0

```

▼ Training the model

```

autoencoder.fit(
    x=train_data,
    y=train_data,
    epochs=50,
    batch_size=128,
    shuffle=True,
    validation_data=(test_data, test_data),
)

Epoch 1/50
469/469 [=====] - 16s 11ms/step - loss: 0.1303 - val_loss: 0.0732
Epoch 2/50
469/469 [=====]
Epoch 3/50           - 4s 8ms/step - loss: 0.0718 - val_loss: 0.0699
469/469 [=====]
Epoch 4/50           - 4s 8ms/step - loss: 0.0695 - val_loss: 0.0684
469/469 [=====]
Epoch 5/50           - 4s 9ms/step - loss: 0.0683 - val_loss: 0.0674
469/469 [=====]
Epoch 6/50           - 4s 8ms/step - loss: 0.0675 - val_loss: 0.0666
469/469 [=====]
Epoch 7/50           - 4s 8ms/step - loss: 0.0669 - val_loss: 0.0661
469/469 [=====]
Epoch 8/50           - 4s 9ms/step - loss: 0.0664 - val_loss: 0.0659
469/469 [=====]
Epoch 9/50           - 4s 8ms/step - loss: 0.0661 - val_loss: 0.0654
469/469 [=====]
Epoch 10/50          - 4s 8ms/step - loss: 0.0657 - val_loss: 0.0652
469/469 [=====]
Epoch 11/50          - 4s 9ms/step - loss: 0.0655 - val_loss: 0.0649
469/469 [=====]
Epoch 12/50          - 4s 8ms/step - loss: 0.0652 - val_loss: 0.0647
469/469 [=====]
Epoch 13/50          - 4s 8ms/step - loss: 0.0650 - val_loss: 0.0646
469/469 [=====]
Epoch 14/50          - 4s 9ms/step - loss: 0.0648 - val_loss: 0.0643
469/469 [=====]
Epoch 15/50          - 4s 8ms/step - loss: 0.0647 - val_loss: 0.0642
469/469 [=====]
Epoch 16/50          - 4s 8ms/step - loss: 0.0645 - val_loss: 0.0641
469/469 [=====]
Epoch 17/50          - 4s 9ms/step - loss: 0.0644 - val_loss: 0.0640
469/469 [=====]
Epoch 18/50          - 4s 8ms/step - loss: 0.0642 - val_loss: 0.0638
469/469 [=====]
Epoch 19/50          - 4s 8ms/step - loss: 0.0641 - val_loss: 0.0637
469/469 [=====]
Epoch 20/50          - 4s 9ms/step - loss: 0.0640 - val_loss: 0.0636
469/469 [=====]
Epoch 21/50          - 4s 8ms/step - loss: 0.0639 - val_loss: 0.0636
469/469 [=====]
Epoch 22/50          - 4s 8ms/step - loss: 0.0638 - val_loss: 0.0634
469/469 [=====]
Epoch 23/50          - 4s 9ms/step - loss: 0.0637 - val_loss: 0.0633
469/469 [=====]
Epoch 24/50          - 4s 8ms/step - loss: 0.0636 - val_loss: 0.0634
469/469 [=====]
Epoch 25/50          - 4s 8ms/step - loss: 0.0636 - val_loss: 0.0633
469/469 [=====]
Epoch 26/50          - 4s 9ms/step - loss: 0.0635 - val_loss: 0.0631
469/469 [=====]
- 4s 8ms/step - loss: 0.0634 - val_loss: 0.0630

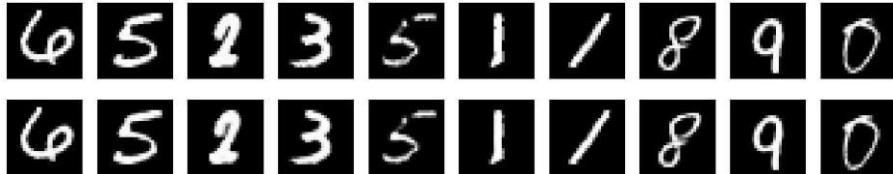
```

```
Epoch 27/50
469/469 [=====] - 4s 8ms/step - loss: 0.0634 - val_loss: 0.0630
Epoch 28/50
469/469 [=====] - 4s 9ms/step - loss: 0.0633 - val_loss: 0.0629
Epoch 29/50
469/469 [=====] - 4s 8ms/step - loss: 0.0632 - val_loss: 0.0628
```

▼ Prediction

```
predictions = autoencoder.predict(test_data)
display(test_data, predictions)
```

```
313/313 [=====] - 1s 2ms/step
```



```
autoencoder.fit(
    x=noisy_train_data,
    y=train_data,
    epochs=100,
    batch_size=128,
    shuffle=True,
    validation_data=(noisy_test_data, test_data),
)
```

```
Epoch 1/100
469/469 [=====] - 4s 9ms/step - loss: 0.1023 - val_loss: 0.0943
Epoch 2/100
469/469 [=====] - 4s 8ms/step - loss: 0.0937 - val_loss: 0.0919
Epoch 3/100
469/469 [=====] - 4s 8ms/step - loss: 0.0920 - val_loss: 0.0906
Epoch 4/100
469/469 [=====] - 4s 9ms/step - loss: 0.0908 - val_loss: 0.0896
Epoch 5/100
469/469 [=====] - 4s 8ms/step - loss: 0.0900 - val_loss: 0.0891
Epoch 6/100
469/469 [=====] - 4s 8ms/step - loss: 0.0894 - val_loss: 0.0884
Epoch 7/100
469/469 [=====] - 4s 9ms/step - loss: 0.0889 - val_loss: 0.0880
Epoch 8/100
469/469 [=====] - 4s 8ms/step - loss: 0.0885 - val_loss: 0.0877
Epoch 9/100
469/469 [=====] - 4s 8ms/step - loss: 0.0881 - val_loss: 0.0874
Epoch 10/100
469/469 [=====] - 4s 9ms/step - loss: 0.0878 - val_loss: 0.0871
Epoch 11/100
469/469 [=====] - 4s 8ms/step - loss: 0.0875 - val_loss: 0.0868
Epoch 12/100
469/469 [=====] - 4s 8ms/step - loss: 0.0873 - val_loss: 0.0866
Epoch 13/100
469/469 [=====] - 4s 9ms/step - loss: 0.0871 - val_loss: 0.0866
Epoch 14/100
469/469 [=====] - 4s 8ms/step - loss: 0.0869 - val_loss: 0.0863
Epoch 15/100
469/469 [=====] - 4s 8ms/step - loss: 0.0868 - val_loss: 0.0863
Epoch 16/100
469/469 [=====] - 4s 9ms/step - loss: 0.0866 - val_loss: 0.0861
Epoch 17/100
469/469 [=====] - 4s 8ms/step - loss: 0.0865 - val_loss: 0.0860
Epoch 18/100
469/469 [=====] - 4s 8ms/step - loss: 0.0864 - val_loss: 0.0859
Epoch 19/100
469/469 [=====] - 4s 9ms/step - loss: 0.0863 - val_loss: 0.0858
Epoch 20/100
469/469 [=====] - 4s 8ms/step - loss: 0.0862 - val_loss: 0.0857
Epoch 21/100
469/469 [=====] - 4s 8ms/step - loss: 0.0861 - val_loss: 0.0857
Epoch 22/100
```

```
469/469 [=====] - 4s 9ms/step - loss: 0.0860 - val_loss: 0.0855
Epoch 23/100
469/469 [=====] - 4s 8ms/step - loss: 0.0859 - val_loss: 0.0854
Epoch 24/100
469/469 [=====] - 4s 8ms/step - loss: 0.0858 - val_loss: 0.0855
Epoch 25/100
469/469 [=====] - 4s 9ms/step - loss: 0.0858 - val_loss: 0.0856
Epoch 26/100
469/469 [=====] - 4s 8ms/step - loss: 0.0857 - val_loss: 0.0852
Epoch 27/100
469/469 [=====] - 4s 8ms/step - loss: 0.0856 - val_loss: 0.0852
Epoch 28/100
469/469 [=====] - 4s 9ms/step - loss: 0.0856 - val_loss: 0.0851
Epoch 29/100
469/469 [=====] - 4s 8ms/step - loss: 0.0855 - val_loss: 0.0851
```

```
predictions = autoencoder.predict(noisy_test_data)
display(noisy_test_data, predictions)
```

